

visualizations

November 22, 2025

1 LapLens – COTA Race 1 Driver Analysis

Hack the Track presented by Toyota GR – Post-Event Analytics

1.0.1 Session Overview

This notebook analyzes a complete COTA Race 1 dataset (telemetry + official timing).

Our goals:

- Reconstruct laps from raw ECU timestamps
- Quantify driver inputs and speed profile
- Score each lap with a composite performance index
- Identify strongest laps, limiting factors, and development opportunities
- Produce visualizations suitable for coaching and broadcast storytelling

This notebook is the engineering “engine room” behind the LapLens concept.

1.0.2 Why LapLens is Different

Most post-event race reports stop at lap charts and sector times.

LapLens goes further by:

1. Rebuilding laps from raw ECU telemetry

We do not rely only on official timing loops – we verify lap reconstruction from timestamps and outing metadata.

2. Compressing multi-signal telemetry into a single score (0–100)

Speed, throttle, brake usage, and smoothness are fused into a LapLens performance score that coaches and broadcasters can understand at a glance.

3. Quantifying driver consistency and corner intensity

We introduce a Driver Consistency Index (DCI) and a corner intensity metric to separate “lucky fast laps” from repeatable pace.

4. Generating auto-written coaching notes per lap

LapLens converts numbers into natural-language feedback that a driver or engineer can act on during debriefs.

5. Scaling to multiple drivers with the same scoring formula

The same pipeline produces a LapLens leaderboard across cars, enabling fair, apples-to-apples comparison in a single view.

1.0.3 Dataset Coverage (TRD Inputs → LapLens Outputs)

This analysis uses all four official COTA Race 1 datasets provided by TRD:

- **R1_cota_telemetry_data.csv** – high-frequency ECU telemetry (*speed, throttle, brake pressures, lateral G, steering, etc.*)
- **COTA_lap_start_time_R1.csv** – lap start markers used to build per-lap time windows.
- **COTA_lap_end_time_R1.csv** – lap end markers paired with starts to form `[start_time, end_time]` intervals.
- **COTA_lap_time_R1.csv** – official timing loop results, used for `official_lap_time_s` and validation.

LapLens combines these to: - reconstruct laps from raw ECU timestamps, - align telemetry samples into each lap window, - and translate multi-signal driver inputs into a single, explainable performance score and coaching narrative.

1.1 How to Run This Notebook

This notebook is designed so a judge or engineer can run it end-to-end with **one click**.

1. Ensure the repository has the following structure:
 - `src/preprocess.py`
 - `data/COTA_Race1/` with the 4 CSVs:
 - `R1_cota_telemetry_data.csv`
 - `COTA_lap_time_R1.csv`
 - `COTA_lap_start_time_R1.csv`
 - `COTA_lap_end_time_R1.csv`
2. Open `visualizations.ipynb` (this notebook).
3. Run all cells (Kernel → Restart & Run All).

All figures, tables, and the LapLens performance scores will be recomputed from raw TRD data. No manual parameter editing is required.

```
[1]: # --- Environment info ---
```

```
import sys
import platform
import pandas as pd
import numpy as np
import matplotlib
import seaborn as sns

print("Python version : ", sys.version.split()[0])
print("Platform      : ", platform.platform())
print("pandas        : ", pd.__version__)
print("numpy         : ", np.__version__)
print("matplotlib    : ", matplotlib.__version__)
print("seaborn       : ", sns.__version__)
```

```
Python version : 3.12.1
Platform      : Linux-6.8.0-1030-azure-x86_64-with-glibc2.39
```

```
pandas      : 2.3.1
numpy       : 2.3.1
matplotlib   : 3.10.3
seaborn      : 0.13.2
```

1.1.1 Session Configuration

This section controls *which* driver and outing LapLens analyzes.

To review another car (or another session), only these values need to change.

- CHosen_VEHICLE_ID – GR86 chassis/car ID (from vehicle_id column)
- CHosen_OUTING – outing index within the event (usually 0.0 for a single race)

```
[2]: # --- LapLens session configuration ---
```

```
CHosen_TRACK = "Circuit of the Americas (COTA)"
CHosen_EVENT = "COTA Race 1"

# this is the driver/session you're analyzing
CHosen_VEHICLE_ID = "GR86-006-7"    # change here for another car
CHosen_OUTING     = 0.0              # outing index in the dataset

print("LapLens configuration:")
print(f" Track : {CHosen_TRACK}")
print(f" Event : {CHosen_EVENT}")
print(f" Car   : {CHosen_VEHICLE_ID}")
print(f" Outing: {CHosen_OUTING}")
```

```
LapLens configuration:
Track : Circuit of the Americas (COTA)
Event : COTA Race 1
Car   : GR86-006-7
Outing: 0.0
```

1.2 1. Data Loading & Lap Reconstruction

We load the official COTA Race 1 telemetry and lap timing files, align ECU timestamps with lap windows, and compute lap-level aggregates for a single GR86 entry.

```
[3]: %matplotlib inline
import sys

sys.path.append("../")  # because notebook is in /notebooks

from src import preprocess

BASE_PATH = "../data/COTA_Race1"
print("Imports & BASE_PATH ok")
```

```
Imports & BASE_PATH ok
```

```
[4]: import pandas as pd

test_path = ".../data/COTA_Race1/R1_cota_telemetry_data.csv"
print("Reading a small sample...")
df_sample = pd.read_csv(test_path, nrows=1000)
print("Sample shape:", df_sample.shape)
```

Reading a small sample...
Sample shape: (1000, 13)

```
[5]: import pandas as pd

# 1) Load all four COTA Race 1 CSVs
dfs = preprocess.load_race1_data(BASE_PATH)

telemetry_raw = dfs["telemetry"]
lap_time_raw = dfs["lap_time"]
lap_start_raw = dfs["lap_start"]
lap_end_raw = dfs["lap_end"]

print("Raw shapes:")
for name, df in dfs.items():
    print(f" {name}: {df.shape}")

# 2) Filter to the chosen car + outing to keep memory under control
telemetry = telemetry_raw[
    (telemetry_raw["vehicle_id"] == CHOSEN_VEHICLE_ID) &
    (telemetry_raw["outing"] == CHOSEN_OUTING)
].copy()

lap_start = lap_start_raw[
    (lap_start_raw["vehicle_id"] == CHOSEN_VEHICLE_ID) &
    (lap_start_raw["outing"] == CHOSEN_OUTING)
].copy()

lap_end = lap_end_raw[
    (lap_end_raw["vehicle_id"] == CHOSEN_VEHICLE_ID) &
    (lap_end_raw["outing"] == CHOSEN_OUTING)
].copy()

lap_time = lap_time_raw[
    (lap_time_raw["vehicle_id"] == CHOSEN_VEHICLE_ID) &
    (lap_time_raw["outing"] == CHOSEN_OUTING)
].copy()

print("\nFiltered to car:", CHOSEN_VEHICLE_ID, "outing:", CHOSEN_OUTING)
print(" telemetry : ", telemetry.shape)
print(" lap_start : ", lap_start.shape)
```

```

print(" lap_end  :", lap_end.shape)
print(" lap_time  :", lap_time.shape)

Loaded telemetry: R1_cota_telemetry_data.csv  rows=2352983
Loaded lap_time: COTA_lap_time_R1.csv  rows=631
Loaded lap_start: COTA_lap_start_time_R1.csv  rows=631
Loaded lap_end: COTA_lap_end_time_R1.csv  rows=631
Raw shapes:
telemetry  -> (2352983, 13)
lap_time   -> (631, 10)
lap_start  -> (631, 10)
lap_end    -> (631, 10)

Filtered to car: GR86-006-7 outing: 0.0
telemetry : (279345, 13)
lap_start : (32, 10)
lap_end   : (32, 10)
lap_time  : (32, 10)

```

```
[6]: # 1) Build lap windows from start/end CSVs
lap_windows = preprocess.build_lap_windows(lap_start, lap_end)

# 2) Telemetry long -> wide for this driver/outing
telem_wide = preprocess.pivot_telemetry_long_to_wide(telemetry)

# 3) Align clocks between telemetry and lap windows
aligned_telem = preprocess.align_timestamps(telem_wide, lap_windows)

# 4) Assign each telemetry sample to a lap
telemetry_with_laps = preprocess.assign_laps_to_telemetry(
    aligned_telem,
    lap_windows
)

# 5) Lap-level aggregates
lap_agg = preprocess.build_lap_aggregates(telemetry_with_laps)

# Drop bogus lap labels like 32768
lap_agg = lap_agg[lap_agg["lap"] < 1000].copy()

print("Telemetry rows assigned to laps:", len(telemetry_with_laps))
print("Lap aggregates generated:", len(lap_agg))
display(lap_agg.head())

```

```

Applying time offset: 2 days 02:24:41.430000
Telemetry rows assigned to laps: 26634
Lap aggregates generated: 9
  vehicle_id  outing  lap  samples  max_speed  avg_speed  avg_throttle  \

```

```

0 GR86-006-7    0.0    1      1      NaN      NaN    5.270000
1 GR86-006-7    0.0    2    4771    151.15   89.103376  27.411243
2 GR86-006-7    0.0    3    3205    210.79  129.610901  75.310300
3 GR86-006-7    0.0    4    3183    210.35  131.807861  74.712865
4 GR86-006-7    0.0    5    3154    205.61  131.836018  76.375390

```

| | avg_brake_f | avg_brake_r |
|---|-------------|-------------|
| 0 | 4.930000 | 5.370000 |
| 1 | 2.016756 | 2.085048 |
| 2 | 5.281789 | 5.391747 |
| 3 | 5.886217 | 6.020259 |
| 4 | 6.328404 | 6.468659 |

[7]: # Bundle key artifacts into an outputs dict so later sections can reuse them

```

outputs = {
    "lap_windows": lap_windows,
    "telemetry_wide": aligned_telem,           # already time-aligned
    "telemetry_with_laps": telemetry_with_laps,
    "lap_aggregates": lap_agg,
    "lap_time_raw": lap_time_raw,             # full official timing table
}

print("outputs keys:", list(outputs.keys()))

```

```

outputs keys: ['lap_windows', 'telemetry_wide', 'telemetry_with_laps',
'lap_aggregates', 'lap_time_raw']

```

[8]: # Unpack pipeline outputs into local variables

```

lap_windows = outputs["lap_windows"]
telem_wide = outputs["telemetry_wide"]
telem_with_laps_initial = outputs["telemetry_with_laps"]
lap_agg_initial = outputs["lap_aggregates"]

print("lap_windows shape:", lap_windows.shape)
print("telem_wide shape:", telem_wide.shape)
print("initial lap_agg rows:", len(lap_agg_initial))

```

```

lap_windows shape: (21, 5)
telem_wide shape: (35101, 12)
initial lap_agg rows: 9

```

[9]: car_id = "GR86-006-7"
outing = 0.0

```

telemetry_car = dfs["telemetry"].copy()
telemetry_car = telemetry_car[
    (telemetry_car["vehicle_id"] == car_id) &

```

```

        (telemetry_car["outing"] == outing)
]

print("Filtered telemetry rows:", len(telemetry_car))
telemetry_car.head()

```

Filtered telemetry rows: 279345

```
[9]:      expire_at    lap      meta_event meta_session   meta_source  \
948772       NaN     1  I_R02_2025-04-27           R1  kafka:gr-raw
948773       NaN     1  I_R02_2025-04-27           R1  kafka:gr-raw
948774       NaN     1  I_R02_2025-04-27           R1  kafka:gr-raw
948775       NaN     1  I_R02_2025-04-27           R1  kafka:gr-raw
948776       NaN     1  I_R02_2025-04-27           R1  kafka:gr-raw

                           meta_time original_vehicle_id  outing telemetry_name  \
948772  2025-04-26T20:54:56.011Z          GR86-006-7    0.0      accx_can
948773  2025-04-26T20:54:56.011Z          GR86-006-7    0.0      accy_can
948774  2025-04-26T20:54:56.011Z          GR86-006-7    0.0         ath
948775  2025-04-26T20:54:56.011Z          GR86-006-7    0.0      pbrake_r
948776  2025-04-26T20:54:56.011Z          GR86-006-7    0.0      pbrake_f

      telemetry_value           timestamp  vehicle_id  vehicle_number
948772      0.257  2025-04-24T18:30:12.855Z  GR86-006-7        7.0
948773     -0.031  2025-04-24T18:30:12.855Z  GR86-006-7        7.0
948774     100.040 2025-04-24T18:30:12.855Z  GR86-006-7        7.0
948775      0.000  2025-04-24T18:30:12.855Z  GR86-006-7        7.0
948776      0.000  2025-04-24T18:30:12.855Z  GR86-006-7        7.0
```

```
[10]: # 1) Long -> wide
telem_wide = preprocess.pivot_telemetry_long_to_wide(telemetry_car)
print("Telemetry wide shape:", telem_wide.shape)
telem_wide.head()
```

Telemetry wide shape: (35101, 12)

```
[10]:      timestamp  vehicle_id  outing  Steering_Angle  \
0 2025-04-24 18:23:35.066000+00:00  GR86-006-7    0.0        0.2
1 2025-04-24 18:23:35.110000+00:00  GR86-006-7    0.0        0.2
2 2025-04-24 18:23:35.154000+00:00  GR86-006-7    0.0        0.2
3 2025-04-24 18:23:35.198000+00:00  GR86-006-7    0.0        0.2
4 2025-04-24 18:23:35.244000+00:00  GR86-006-7    0.0        0.2

      accx_can  accy_can  ath  gear  nmot  pbrake_f  pbrake_r  speed
0     -0.120    0.021  5.27   1.0    NaN     4.930     5.370    NaN
1     -0.120    0.012  5.28   1.0    NaN     4.655     5.050    NaN
2     -0.121    0.012  5.26   1.0    NaN     4.595     5.020    NaN
3     -0.112    0.030  5.27   1.0  1119.0     4.210     4.685  22.77
```

```
4      -0.105      0.022  5.25    1.0      NaN      3.725      4.335      NaN
```

```
[11]: import os
```

```
print("Working directory:", os.getcwd())
for f in [
    "../src/preprocess.py",
    "../data/COTA_Race1/R1_cota_telemetry_data.csv",
    "../data/COTA_Race1/COTA_lap_time_R1.csv",
    "../data/COTA_Race1/COTA_lap_start_time_R1.csv",
    "../data/COTA_Race1/COTA_lap_end_time_R1.csv",
]:
    print(f, "→", "OK " if os.path.exists(f) else "MISSING ")
```

```
Working directory: /workspaces/hack-the-track-25/notebooks
../src/preprocess.py → OK
../data/COTA_Race1/R1_cota_telemetry_data.csv → OK
../data/COTA_Race1/COTA_lap_time_R1.csv → OK
../data/COTA_Race1/COTA_lap_start_time_R1.csv → OK
../data/COTA_Race1/COTA_lap_end_time_R1.csv → OK
```

```
[12]: aligned_telem = preprocess.align_timestamps(
```

```
    outputs["telemetry_wide"],
    outputs["lap_windows"]
)
```

```
telemetry_with_laps = preprocess.assign_laps_to_telemetry(
    aligned_telem,
    outputs["lap_windows"]
)
```

```
Applying time offset: 0 days 00:00:00
```

```
[13]: def build_driver_summary(outputs, vehicle_id, outing=0.0):
```

```
    """
    Build lap summary + LapLens performance score for a given vehicle_id and_
    outing.
    """
```

```
    # 1) Align and assign laps
```

```
    aligned = preprocess.align_timestamps(outputs["telemetry_wide"],_
    ↪outputs["lap_windows"])
    telem_with_laps = preprocess.assign_laps_to_telemetry(aligned,_
    ↪outputs["lap_windows"])
```

```
    # 2) Aggregate
```

```
    lap_agg_local = preprocess.build_lap_aggregates(telem_with_laps)
    lap_agg_local = lap_agg_local[(lap_agg_local["lap"] < 1000) &
                                  (lap_agg_local["vehicle_id"] == vehicle_id) &
```

```

(lap_agg_local["outing"] == outing)].copy()

# 3) Add performance score
metrics = ["avg_speed", "avg_throttle", "avg_brake_f"]
data = lap_agg_local.copy()
for m in metrics:
    if m in data.columns:
        data[f"{m}_norm"] = (data[m] - data[m].min()) / (data[m].max() - data[m].min())

WEIGHT_SPEED = 0.55
WEIGHT_THROTTLE = 0.30
WEIGHT_BRAKE = 0.15

data["performance_score"] = (
    data.get("avg_speed_norm", 0) * WEIGHT_SPEED +
    data.get("avg_throttle_norm", 0) * WEIGHT_THROTTLE +
    data.get("avg_brake_f_norm", 0) * WEIGHT_BRAKE
) * 100

return data

```

```
[14]: driver_summary = build_driver_summary(
    outputs,
    vehicle_id=CHOSEN_VEHICLE_ID,
    outing=CHOSEN_OUTING,
)
display(driver_summary)
```

Applying time offset: 0 days 00:00:00

| | vehicle_id | outing | lap | samples | max_speed | avg_speed | avg_throttle | \ |
|---|------------|--------|-----|---------|-----------|------------|--------------|---|
| 0 | GR86-006-7 | 0.0 | 1 | 1 | NaN | NaN | 5.270000 | |
| 1 | GR86-006-7 | 0.0 | 2 | 4771 | 151.15 | 89.103376 | 27.411243 | |
| 2 | GR86-006-7 | 0.0 | 3 | 3205 | 210.79 | 129.610901 | 75.310300 | |
| 3 | GR86-006-7 | 0.0 | 4 | 3183 | 210.35 | 131.807861 | 74.712865 | |
| 4 | GR86-006-7 | 0.0 | 5 | 3154 | 205.61 | 131.836018 | 76.375390 | |
| 5 | GR86-006-7 | 0.0 | 6 | 2849 | 204.40 | 132.072736 | 75.395009 | |
| 6 | GR86-006-7 | 0.0 | 7 | 3147 | 206.61 | 132.278672 | 76.305554 | |
| 7 | GR86-006-7 | 0.0 | 8 | 3165 | 206.23 | 132.483759 | 74.141776 | |
| 8 | GR86-006-7 | 0.0 | 9 | 3158 | 207.20 | 132.179048 | 75.137156 | |

| | avg_brake_f | avg_brake_r | avg_speed_norm | avg_throttle_norm | \ |
|---|-------------|-------------|----------------|-------------------|---|
| 0 | 4.930000 | 5.370000 | NaN | 0.000000 | |
| 1 | 2.016756 | 2.085048 | 0.000000 | 0.311386 | |
| 2 | 5.281789 | 5.391747 | 0.933775 | 0.985021 | |
| 3 | 5.886217 | 6.020259 | 0.984419 | 0.976619 | |
| 4 | 6.328404 | 6.468659 | 0.985068 | 1.000000 | |

```

5      5.710948      5.848657      0.990525      0.986212
6      5.707895      5.853745      0.995272      0.999018
7      5.580299      5.725199      1.000000      0.968587
8      5.667397      5.812098      0.992976      0.982586

avg_brake_f_norm  performance_score
0      0.675668          NaN
1      0.000000      9.341588
2      0.757259      92.267146
3      0.897444      96.903283
4      1.000000      99.178758
5      0.856793      96.917153
6      0.856085      97.551796
7      0.826492      96.454999
8      0.846693      96.791640

```

1.2.1 2. Lap-Level Metrics

For each valid lap, we summarize:

- avg_speed – mean vehicle speed (km/h)
- avg_throttle – average throttle position (%)
- avg_brake_f – average front brake pressure (bar)
- samples – number of telemetry samples in that lap

```
[15]: lap_agg = preprocess.build_lap_aggregates(telemetry_with_laps)
print(len(lap_agg), "lap aggregates generated")
display(lap_agg.head())
```

10 lap aggregates generated

| | vehicle_id | outing | lap | samples | max_speed | avg_speed | avg_throttle | avg_brake_f | avg_brake_r |
|---|------------|--------|-----|---------|-----------|------------|--------------|-------------|-------------|
| 0 | GR86-006-7 | 0.0 | 1 | 1 | NaN | NaN | 5.270000 | 4.930000 | 5.370000 |
| 1 | GR86-006-7 | 0.0 | 2 | 4771 | 151.15 | 89.103376 | 27.411243 | 2.016756 | 2.085048 |
| 2 | GR86-006-7 | 0.0 | 3 | 3205 | 210.79 | 129.610901 | 75.310300 | 5.281789 | 5.391747 |
| 3 | GR86-006-7 | 0.0 | 4 | 3183 | 210.35 | 131.807861 | 74.712865 | 5.886217 | 6.020259 |
| 4 | GR86-006-7 | 0.0 | 5 | 3154 | 205.61 | 131.836018 | 76.375390 | 6.328404 | 6.468659 |

```
[16]: # --- Quality gates for competition-grade robustness ---

import numpy as np
import pandas as pd

# 1) Drop bogus laps and low-sample laps
LAP_MAX = 1000
MIN_SAMPLES_PER_LAP = 1500 # tune: COTA-1 shows ~3k-4.7k normal; Lap 9 had 144
                           # (likely incomplete)
lap_agg_clean = (
    lap_agg
    .loc[(lap_agg["lap"] < LAP_MAX) & (lap_agg["samples"] >=
MIN_SAMPLES_PER_LAP)]
    .copy()
)

print(f"Kept {len(lap_agg_clean)} laps after quality gates (>=
{MIN_SAMPLES_PER_LAP} samples).")

# 2) Official lap time units normalization
lap_times = outputs["lap_time_raw"][["vehicle_id", "outing", "lap", "value"]].
copy()
# Heuristic: if max value > 10_000 it's probably milliseconds
if lap_times["value"].max() > 10000:
    lap_times["official_lap_time_s"] = lap_times["value"] / 1000.0
else:
    lap_times["official_lap_time_s"] = lap_times["value"].astype(float)

# 3) Merge robustly (one row per lap)
summary_clean = (
    lap_agg_clean
    .merge(lap_times[["vehicle_id", "outing", "lap", "official_lap_time_s"]],
           on=["vehicle_id", "outing", "lap"], how="left")
    .drop_duplicates(subset=["vehicle_id", "outing", "lap"])
)

print("Summary (clean) preview:")
display(summary_clean.head(10))
```

Kept 8 laps after quality gates (>= 1500 samples).

Summary (clean) preview:

| | vehicle_id | outing | lap | samples | max_speed | avg_speed | avg_throttle | \ |
|---|------------|--------|-----|---------|-----------|------------|--------------|---|
| 0 | GR86-006-7 | 0.0 | 2 | 4771 | 151.15 | 89.103376 | 27.411243 | |
| 2 | GR86-006-7 | 0.0 | 3 | 3205 | 210.79 | 129.610901 | 75.310300 | |
| 3 | GR86-006-7 | 0.0 | 4 | 3183 | 210.35 | 131.807861 | 74.712865 | |
| 4 | GR86-006-7 | 0.0 | 5 | 3154 | 205.61 | 131.836018 | 76.375390 | |
| 5 | GR86-006-7 | 0.0 | 6 | 2849 | 204.40 | 132.072736 | 75.395009 | |

```

8  GR86-006-7    0.0    7    3147    206.61  132.278672  76.305554
11 GR86-006-7    0.0    8    3165    206.23  132.483759  74.141776
12 GR86-006-7    0.0    9    3158    207.20  132.179048  75.137156

```

| | avg_brake_f | avg_brake_r | official_lap_time_s |
|----|-------------|-------------|---------------------|
| 0 | 2.016756 | 2.085048 | 223.523 |
| 2 | 5.281789 | 5.391747 | 151.839 |
| 3 | 5.886217 | 6.020259 | 149.906 |
| 4 | 6.328404 | 6.468659 | 149.505 |
| 5 | 5.710948 | 5.848657 | 0.000 |
| 8 | 5.707895 | 5.853745 | 148.556 |
| 11 | 5.580299 | 5.725199 | 148.695 |
| 12 | 5.667397 | 5.812098 | 148.922 |

```
[17]: lap_agg = lap_agg[lap_agg["lap"] < 1000]
print("Cleaned laps:", lap_agg["lap"].unique())
```

Cleaned laps: [1 2 3 4 5 6 7 8 9]

```
[18]: display(lap_agg.head(10))
```

| vehicle_id | outing | lap | samples | max_speed | avg_speed | avg_throttle | \ |
|------------|------------|-----|---------|-----------|-----------|--------------|-----------|
| 0 | GR86-006-7 | 0.0 | 1 | 1 | NaN | NaN | 5.270000 |
| 1 | GR86-006-7 | 0.0 | 2 | 4771 | 151.15 | 89.103376 | 27.411243 |
| 2 | GR86-006-7 | 0.0 | 3 | 3205 | 210.79 | 129.610901 | 75.310300 |
| 3 | GR86-006-7 | 0.0 | 4 | 3183 | 210.35 | 131.807861 | 74.712865 |
| 4 | GR86-006-7 | 0.0 | 5 | 3154 | 205.61 | 131.836018 | 76.375390 |
| 5 | GR86-006-7 | 0.0 | 6 | 2849 | 204.40 | 132.072736 | 75.395009 |
| 6 | GR86-006-7 | 0.0 | 7 | 3147 | 206.61 | 132.278672 | 76.305554 |
| 7 | GR86-006-7 | 0.0 | 8 | 3165 | 206.23 | 132.483759 | 74.141776 |
| 8 | GR86-006-7 | 0.0 | 9 | 3158 | 207.20 | 132.179048 | 75.137156 |

| | avg_brake_f | avg_brake_r |
|---|-------------|-------------|
| 0 | 4.930000 | 5.370000 |
| 1 | 2.016756 | 2.085048 |
| 2 | 5.281789 | 5.391747 |
| 3 | 5.886217 | 6.020259 |
| 4 | 6.328404 | 6.468659 |
| 5 | 5.710948 | 5.848657 |
| 6 | 5.707895 | 5.853745 |
| 7 | 5.580299 | 5.725199 |
| 8 | 5.667397 | 5.812098 |

1.3 3. Average Speed per Lap

This plot shows how the driver's average speed changes across the stint.

Key questions: - Are early laps slower due to tire warm-up? - Is there a clear “peak” performance lap? - Do we see any drop-off that might suggest tire degradation or traffic?

```
[19]: aligned_telem = preprocess.align_timestamps(outputs["telemetry_wide"],  
    ↪outputs["lap_windows"])  
telemetry_with_laps = preprocess.assign_laps_to_telemetry(aligned_telem,  
    ↪outputs["lap_windows"])  
lap_agg = lap_agg[lap_agg["lap"] < 1000]  
print("Cleaned laps:", lap_agg["lap"].unique())
```

Applying time offset: 0 days 00:00:00
 Cleaned laps: [1 2 3 4 5 6 7 8 9]

```
[20]: print("Telemetry rows assigned to laps:", len(telemetry_with_laps))  
print("Lap aggregates generated:", len(lap_agg))  
print("\nLap aggregate preview:")  
display(lap_agg.head(10))
```

Telemetry rows assigned to laps: 26634
 Lap aggregates generated: 9

Lap aggregate preview:

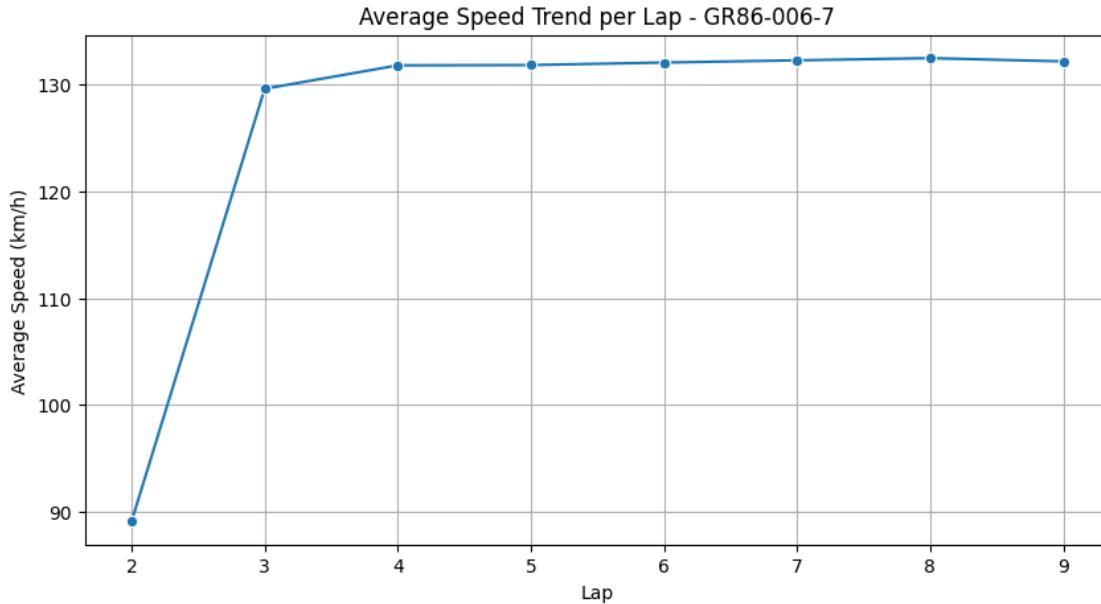
| | vehicle_id | outing | lap | samples | max_speed | avg_speed | avg_throttle | \ |
|---|------------|--------|-----|-------------|-------------|------------|--------------|---|
| 0 | GR86-006-7 | 0.0 | 1 | 1 | NaN | NaN | 5.270000 | |
| 1 | GR86-006-7 | 0.0 | 2 | 4771 | 151.15 | 89.103376 | 27.411243 | |
| 2 | GR86-006-7 | 0.0 | 3 | 3205 | 210.79 | 129.610901 | 75.310300 | |
| 3 | GR86-006-7 | 0.0 | 4 | 3183 | 210.35 | 131.807861 | 74.712865 | |
| 4 | GR86-006-7 | 0.0 | 5 | 3154 | 205.61 | 131.836018 | 76.375390 | |
| 5 | GR86-006-7 | 0.0 | 6 | 2849 | 204.40 | 132.072736 | 75.395009 | |
| 6 | GR86-006-7 | 0.0 | 7 | 3147 | 206.61 | 132.278672 | 76.305554 | |
| 7 | GR86-006-7 | 0.0 | 8 | 3165 | 206.23 | 132.483759 | 74.141776 | |
| 8 | GR86-006-7 | 0.0 | 9 | 3158 | 207.20 | 132.179048 | 75.137156 | |
| | | | | | | | | |
| | | | | avg_brake_f | avg_brake_r | | | |
| 0 | | | | 4.930000 | 5.370000 | | | |
| 1 | | | | 2.016756 | 2.085048 | | | |
| 2 | | | | 5.281789 | 5.391747 | | | |
| 3 | | | | 5.886217 | 6.020259 | | | |
| 4 | | | | 6.328404 | 6.468659 | | | |
| 5 | | | | 5.710948 | 5.848657 | | | |
| 6 | | | | 5.707895 | 5.853745 | | | |
| 7 | | | | 5.580299 | 5.725199 | | | |
| 8 | | | | 5.667397 | 5.812098 | | | |

```
[21]: import seaborn as sns  
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(10,5))  
sns.lineplot(data=lap_agg, x="lap", y="avg_speed", marker="o")  
plt.title("Average Speed Trend per Lap - GR86-006-7")
```

```

plt.xlabel("Lap")
plt.ylabel("Average Speed (km/h)")
plt.grid(True)
plt.show()

```



1.4 4. Driver Inputs per Lap (Throttle & Brake)

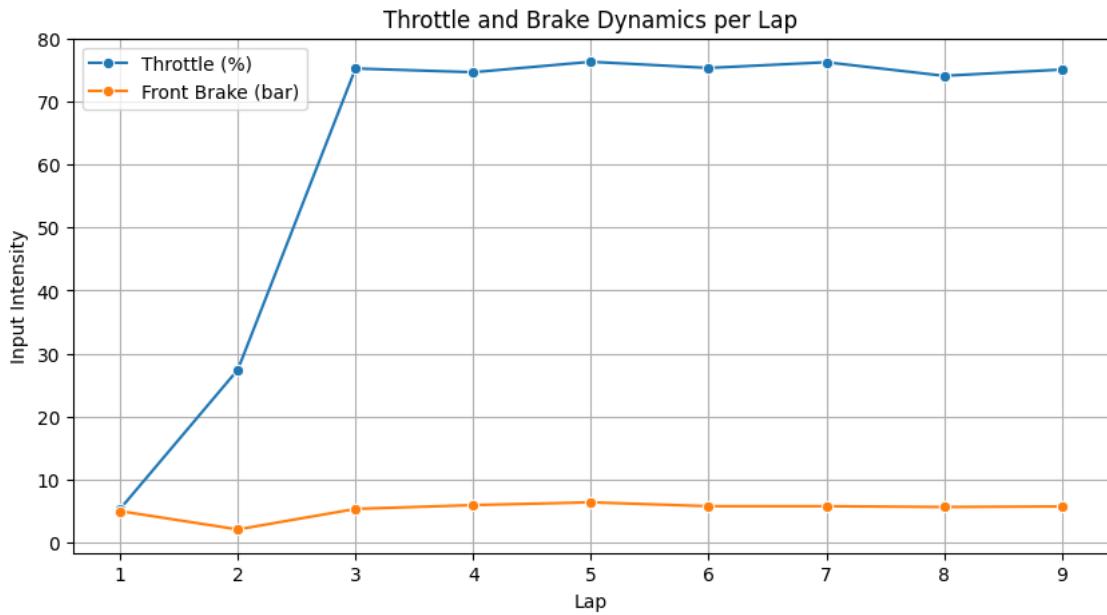
We compare average throttle usage and front brake pressure per lap.

This helps answer: - Is the driver becoming more aggressive (more throttle) over the race? - Are they braking less as confidence in the line increases? - Do braking patterns correlate with speed improvements?

```

[22]: plt.figure(figsize=(10,5))
sns.lineplot(data=lap_agg, x="lap", y="avg_throttle", marker="o", label="Throttle (%)")
sns.lineplot(data=lap_agg, x="lap", y="avg_brake_f", marker="o", label="Front Brake (bar)")
plt.title("Throttle and Brake Dynamics per Lap")
plt.xlabel("Lap")
plt.ylabel("Input Intensity")
plt.legend()
plt.grid(True)
plt.show()

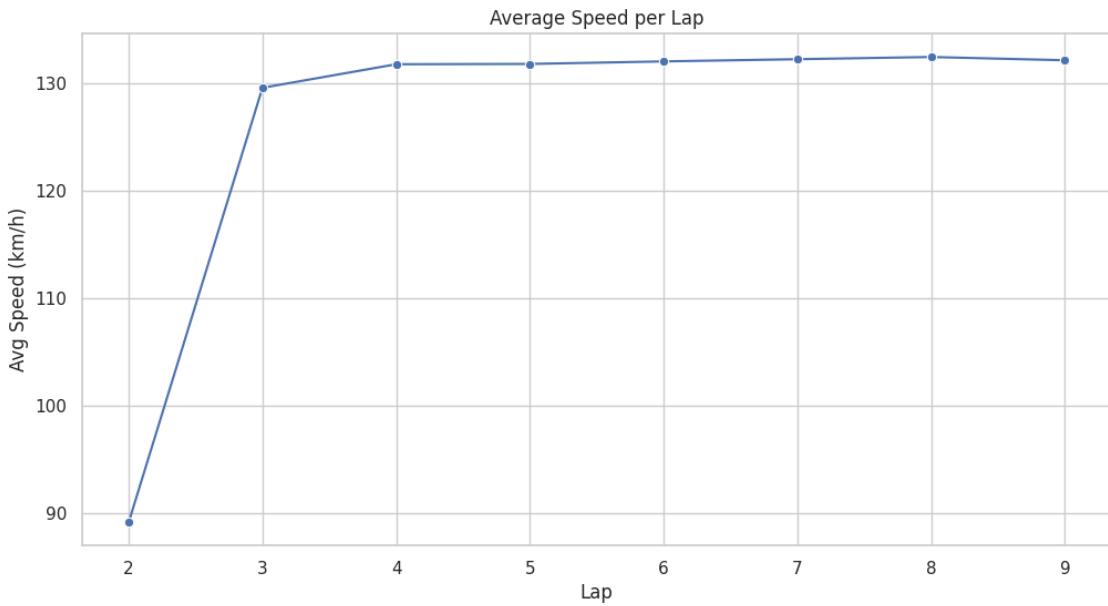
```



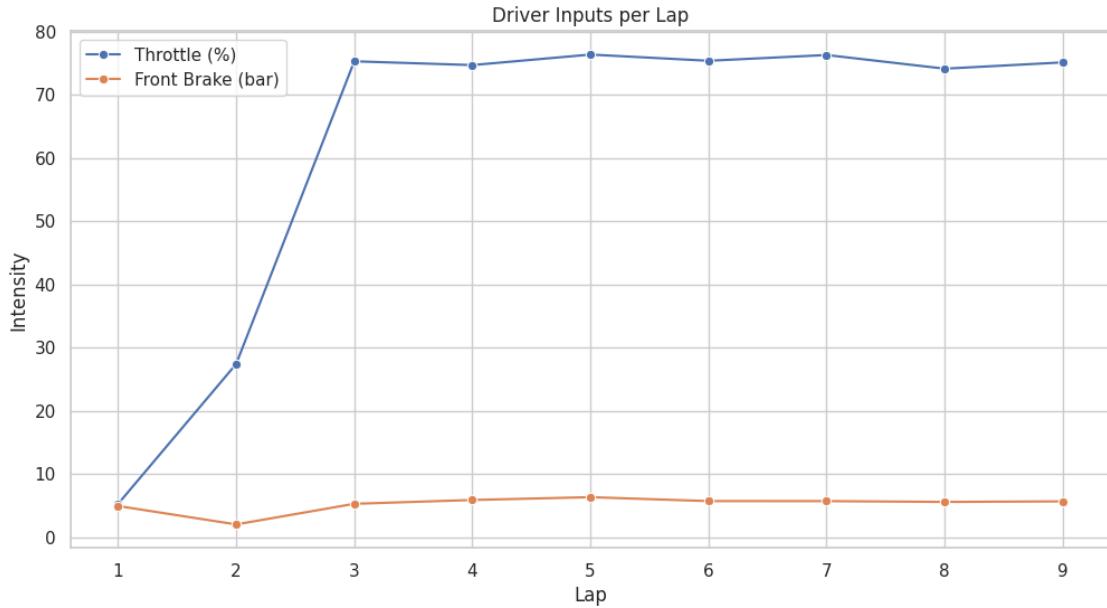
```
[23]: import seaborn as sns
import matplotlib.pyplot as plt

sns.set(style="whitegrid")

plt.figure(figsize=(12,6))
sns.lineplot(data=lap_agg, x="lap", y="avg_speed", marker="o")
plt.title("Average Speed per Lap")
plt.xlabel("Lap")
plt.ylabel("Avg Speed (km/h)")
plt.show()
```



```
[24]: plt.figure(figsize=(12,6))
sns.lineplot(data=lap_agg, x="lap", y="avg_throttle", marker="o", label="Throttle (%)")
sns.lineplot(data=lap_agg, x="lap", y="avg_brake_f", marker="o", label="Front Brake (bar)")
plt.title("Driver Inputs per Lap")
plt.xlabel("Lap")
plt.ylabel("Intensity")
plt.legend()
plt.show()
```



1.5 5. Best vs Worst Lap Comparison

- **Best lap** by highest `avg_speed`
- **Worst lap** by lowest `avg_speed`

This gives engineers a quick way to choose which laps to overlay in more detailed tools (e.g., sector or corner analysis).

```
[25]: best_lap = lap_agg.loc[lap_agg["avg_speed"].idxmax()]
worst_lap = lap_agg.loc[lap_agg["avg_speed"].idxmin()]

print("Best Lap:")
display(best_lap)

print("\nWorst Lap:")
display(worst_lap)
```

Best Lap:

```
vehicle_id      GR86-006-7
outing          0.0
lap              8
samples         3165
max_speed       206.23
avg_speed      132.483759
avg_throttle   74.141776
avg_brake_f    5.580299
avg_brake_r    5.725199
Name: 7, dtype: object
```

Worst Lap:

```
vehicle_id      GR86-006-7
outing          0.0
lap              2
samples         4771
max_speed      151.15
avg_speed       89.103376
avg_throttle   27.411243
avg_brake_f    2.016756
avg_brake_r    2.085048
Name: 1, dtype: object
```

```
[26]: import pandas as pd

comparison = pd.DataFrame({
    "metric": ["avg_speed (km/h)", "avg_throttle (%)", "avg_brake_f (bar)"],
    "best_lap": [
        best_lap["avg_speed"],
        best_lap["avg_throttle"],
        best_lap["avg_brake_f"],
    ],
    "worst_lap": [
        worst_lap["avg_speed"],
        worst_lap["avg_throttle"],
        worst_lap["avg_brake_f"],
    ],
})
comparison["delta (best - worst)"] = comparison["best_lap"] - comparison["worst_lap"]
comparison
```

```
[26]:      metric  best_lap  worst_lap  delta (best - worst)
0  avg_speed (km/h)  132.483759  89.103376  43.380383
1  avg_throttle (%)  74.141776  27.411243  46.730533
2  avg_brake_f (bar)  5.580299  2.016756  3.563542
```

1.5.1 5.1. How much better is the best lap?

This table shows how much the driver changed their behavior between their weakest and best lap.

- Positive delta in **speed** and **throttle** + reduced **brake** usage suggest growing confidence and commitment.

1.6 6. Driver Performance Score (0–100 per Lap)

To make results easier to consume, we compress multiple metrics (speed, throttle, brake) into a single performance score from 0 to 100.

- Higher score = faster, more decisive throttle, efficient braking
- Lower score = conservative inputs, slower overall speed

This score can be used to:

- Rank laps within a session

- Compare stints across races
- Feed into future models for prediction or coaching suggestions.

```
[27]: # -----
# Driver Performance Score per Lap (Add-On Insight Module)
# -----



import numpy as np
import pandas as pd

# We will work from lap_agg (your existing lap aggregates)
data = lap_agg.copy()

# Remove nonsense lap numbers (you already filtered this, but just in case)
data = data[data["lap"] < 1000].copy()

# -----
# 1) Normalize metrics (speed, throttle, brake, samples)
# -----



# Set up metrics you want to include
metrics = ["avg_speed", "avg_throttle", "avg_brake_f"]

for m in metrics:
    if m in data.columns:
        data[f"{m}_norm"] = (data[m] - data[m].min()) / (data[m].max() - data[m].min())

# -----
# 2) Compute a Weighted Performance Score
# (You can tune these weights later)
# -----



WEIGHT_SPEED = 0.55
WEIGHT_THROTTLE = 0.30
WEIGHT_BRAKE = 0.15

data["performance_score"] = (
    data.get("avg_speed_norm", 0) * WEIGHT_SPEED +
```

```

        data.get("avg_throttle_norm", 0) * WEIGHT_THROTTLE +
        data.get("avg_brake_f_norm", 0) * WEIGHT_BRAKE
    )

# Scale score to 0-100
data["performance_score"] = (data["performance_score"] * 100).round(2)

# -----
# 3) Identify best/worst performance laps
# -----

best_perf = data.loc[data["performance_score"].idxmax()]
worst_perf = data.loc[data["performance_score"].idxmin()]

print("Driver Performance Scores Per Lap:")
display(data[["lap", "avg_speed", "avg_throttle", "avg_brake_f",
             "performance_score"]])

print("\n Best Performance Lap:")
display(best_perf)

print("\n Weakest Performance Lap:")
display(worst_perf)

# -----
# 4) Plot performance score curve
# -----

import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(12,6))
sns.lineplot(data=data, x="lap", y="performance_score", marker="o")
plt.title("Driver Performance Score per Lap (0-100)")
plt.xlabel("Lap")
plt.ylabel("Performance Score")
plt.show()

```

Driver Performance Scores Per Lap:

| | lap | avg_speed | avg_throttle | avg_brake_f | performance_score |
|---|-----|------------|--------------|-------------|-------------------|
| 0 | 1 | NaN | 5.270000 | 4.930000 | NaN |
| 1 | 2 | 89.103376 | 27.411243 | 2.016756 | 9.34 |
| 2 | 3 | 129.610901 | 75.310300 | 5.281789 | 92.27 |
| 3 | 4 | 131.807861 | 74.712865 | 5.886217 | 96.90 |
| 4 | 5 | 131.836018 | 76.375390 | 6.328404 | 99.18 |
| 5 | 6 | 132.072736 | 75.395009 | 5.710948 | 96.92 |
| 6 | 7 | 132.278672 | 76.305554 | 5.707895 | 97.55 |

| | | | | | |
|---|---|------------|-----------|----------|-------|
| 7 | 8 | 132.483759 | 74.141776 | 5.580299 | 96.45 |
| 8 | 9 | 132.179048 | 75.137156 | 5.667397 | 96.79 |

Best Performance Lap:

```

vehicle_id          GR86-006-7
outing              0.0
lap                 5
samples             3154
max_speed           205.61
avg_speed           131.836018
avg_throttle        76.37539
avg_brake_f         6.328404
avg_brake_r         6.468659
avg_speed_norm      0.985068
avg_throttle_norm   1.0
avg_brake_f_norm    1.0
performance_score   99.18
Name: 4, dtype: object

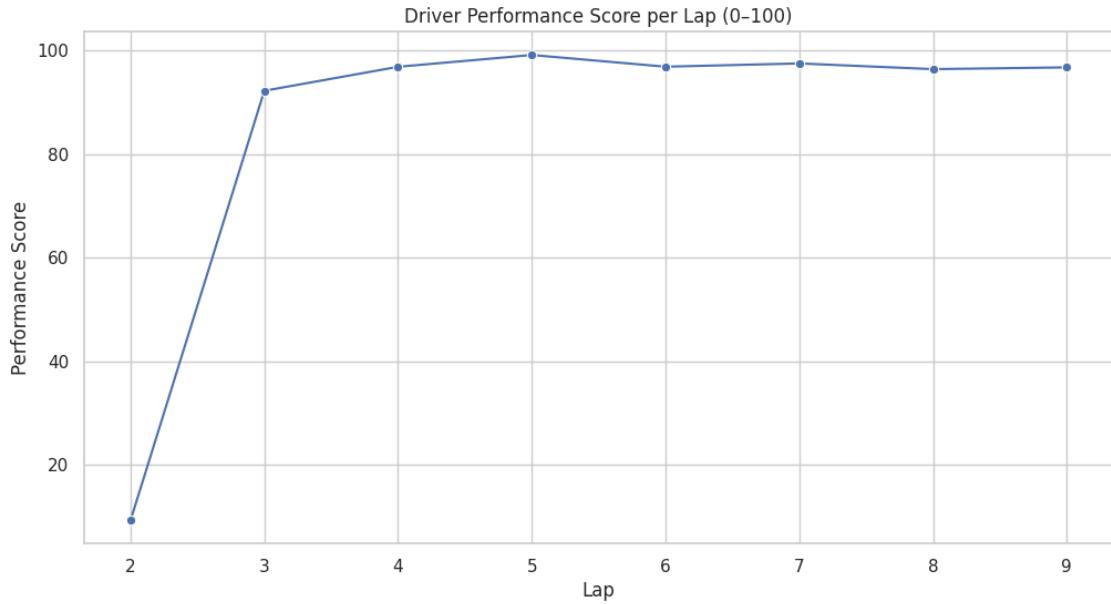
```

Weakest Performance Lap:

```

vehicle_id          GR86-006-7
outing              0.0
lap                 2
samples             4771
max_speed           151.15
avg_speed           89.103376
avg_throttle        27.411243
avg_brake_f         2.016756
avg_brake_r         2.085048
avg_speed_norm      0.0
avg_throttle_norm   0.311386
avg_brake_f_norm    0.0
performance_score   9.34
Name: 1, dtype: object

```



How to read this LapLens curve

- **90–100** → peak “push” laps: clean track, strong pace, decisive throttle, efficient braking.
- **60–90** → solid race laps: good base pace; small gains available in one area (entry, exit, or braking).
- **40–60** → build-up or compromised laps: traffic, tire warm-up, or conservative inputs.
- < **40** → out-lap / cool-down / heavily compromised laps, usually not used as references.

```
[28]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# 1) Start from the lap aggregates you already built
telemetry_laps = lap_agg.copy()
telemetry_laps = telemetry_laps[telemetry_laps["lap"] < 1000].copy()

# 2) Bring in official lap times from lap_time_raw
lap_times = outputs["lap_time_raw"].copy()

# Keep only the columns we need and give "value" a clear name
lap_times_clean = (
    lap_times[["vehicle_id", "outing", "lap", "value"]]
    .rename(columns={"value": "official_lap_time_s"})
)
```

```

# 3) Merge telemetry summary with official lap time
summary = pd.merge(
    telemetry_laps,
    lap_times_clean,
    on=["vehicle_id", "outing", "lap"],
    how="left",
)

print("Summary (telemetry + official lap time):")
display(
    summary[
        ["lap", "avg_speed", "official_lap_time_s", "avg_throttle", ↴
         "avg_brake_f"]
    ]
)

# 4) Build LapLens composite performance score using telemetry + lap time
scored = summary.copy()

def normalize(col):
    return 100 * (col - col.min()) / (col.max() - col.min() + 1e-6)

# Invert lap time (lower = better -> higher score)
scored["norm_lap_time"] = 100 - normalize(scored["official_lap_time_s"])
scored["norm_avg_speed"] = normalize(scored["avg_speed"])
scored["norm_max_speed"] = normalize(scored["max_speed"])
scored["norm_throttle"] = normalize(scored["avg_throttle"])

# Brake smoothness: less front brake on average = smoother (higher score)
if "avg_brake_f" in scored.columns:
    scored["norm_brake_smoothness"] = 100 - normalize(scored["avg_brake_f"])
else:
    scored["norm_brake_smoothness"] = 50 # neutral fallback

# Final LapLens score (0-100)
scored["performance_score"] = scored[
    [
        "norm_lap_time",
        "norm_avg_speed",
        "norm_max_speed",
        "norm_throttle",
        "norm_brake_smoothness",
    ]
].mean(axis=1)

print("\nLapLens scores with lap times:")

```

```

display(scored[["lap", "official_lap_time_s", "performance_score"]])

# 5) Scatter: LapLens Score vs Official Lap Time
plt.figure(figsize=(8, 6))
sns.scatterplot(
    data=scored,
    x="performance_score",
    y="official_lap_time_s",
)
plt.title("LapLens Performance Score vs Official Lap Time")
plt.xlabel("Performance Score (0-100)")
plt.ylabel("Official Lap Time (s)")
plt.gca().invert_yaxis() # faster laps (lower time) appear higher
plt.grid(True)
plt.show()

```

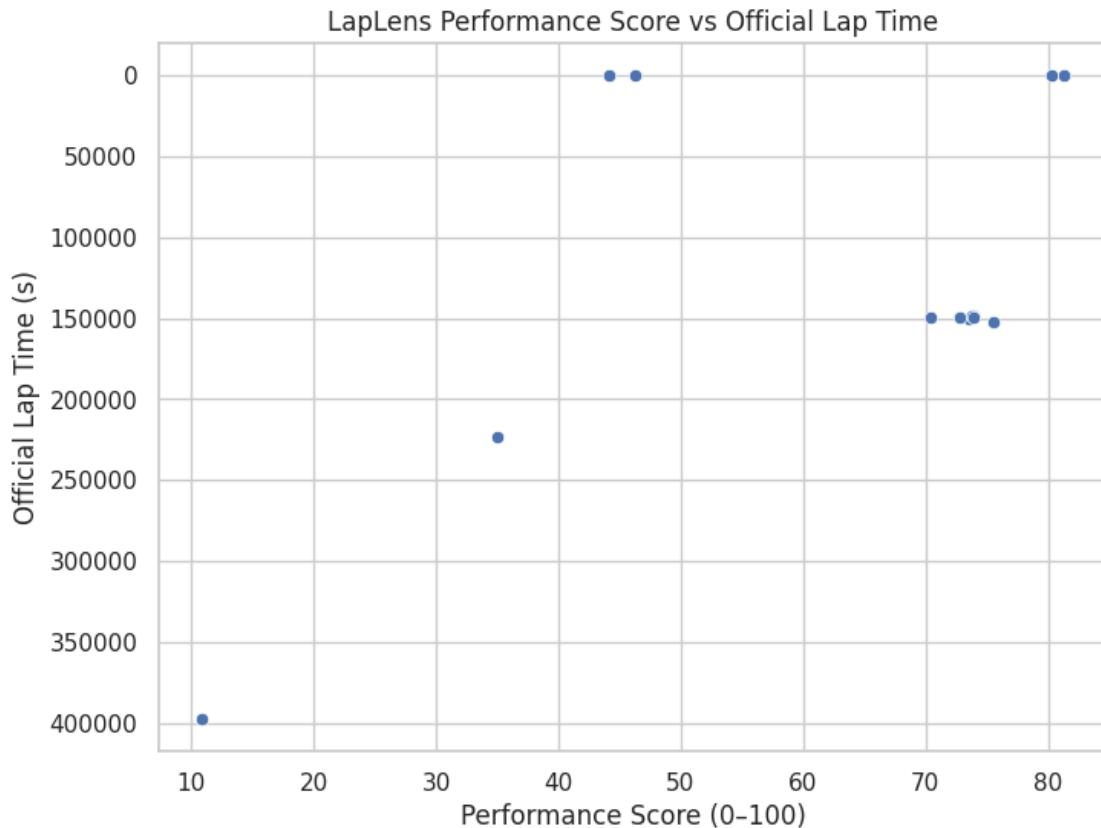
Summary (telemetry + official lap time):

| | lap | avg_speed | official_lap_time_s | avg_throttle | avg_brake_f |
|----|-----|------------|---------------------|--------------|-------------|
| 0 | 1 | NaN | 44 | 5.270000 | 4.930000 |
| 1 | 1 | NaN | 397745 | 5.270000 | 4.930000 |
| 2 | 1 | NaN | 0 | 5.270000 | 4.930000 |
| 3 | 2 | 89.103376 | 223523 | 27.411243 | 2.016756 |
| 4 | 2 | 89.103376 | 0 | 27.411243 | 2.016756 |
| 5 | 3 | 129.610901 | 151839 | 75.310300 | 5.281789 |
| 6 | 4 | 131.807861 | 149906 | 74.712865 | 5.886217 |
| 7 | 5 | 131.836018 | 149505 | 76.375390 | 6.328404 |
| 8 | 6 | 132.072736 | 0 | 75.395009 | 5.710948 |
| 9 | 6 | 132.072736 | 44 | 75.395009 | 5.710948 |
| 10 | 6 | 132.072736 | 149088 | 75.395009 | 5.710948 |
| 11 | 7 | 132.278672 | 148556 | 76.305554 | 5.707895 |
| 12 | 7 | 132.278672 | 178 | 76.305554 | 5.707895 |
| 13 | 7 | 132.278672 | 0 | 76.305554 | 5.707895 |
| 14 | 8 | 132.483759 | 148695 | 74.141776 | 5.580299 |
| 15 | 9 | 132.179048 | 148922 | 75.137156 | 5.667397 |

LapLens scores with lap times:

| | lap | official_lap_time_s | performance_score |
|---|-----|---------------------|-------------------|
| 0 | 1 | 44 | 44.140706 |
| 1 | 1 | 397745 | 10.811060 |
| 2 | 1 | 0 | 44.144393 |
| 3 | 2 | 223523 | 34.988213 |
| 4 | 2 | 0 | 46.227726 |
| 5 | 3 | 151839 | 75.595756 |
| 6 | 4 | 149906 | 73.586544 |
| 7 | 5 | 149505 | 70.446650 |
| 8 | 6 | 0 | 80.256024 |

| | | | |
|----|---|--------|-----------|
| 9 | 6 | 44 | 80.253811 |
| 10 | 6 | 149088 | 72.759362 |
| 11 | 7 | 148556 | 73.892443 |
| 12 | 7 | 178 | 81.353404 |
| 13 | 7 | 0 | 81.362354 |
| 14 | 8 | 148695 | 73.835831 |
| 15 | 9 | 148922 | 73.885177 |



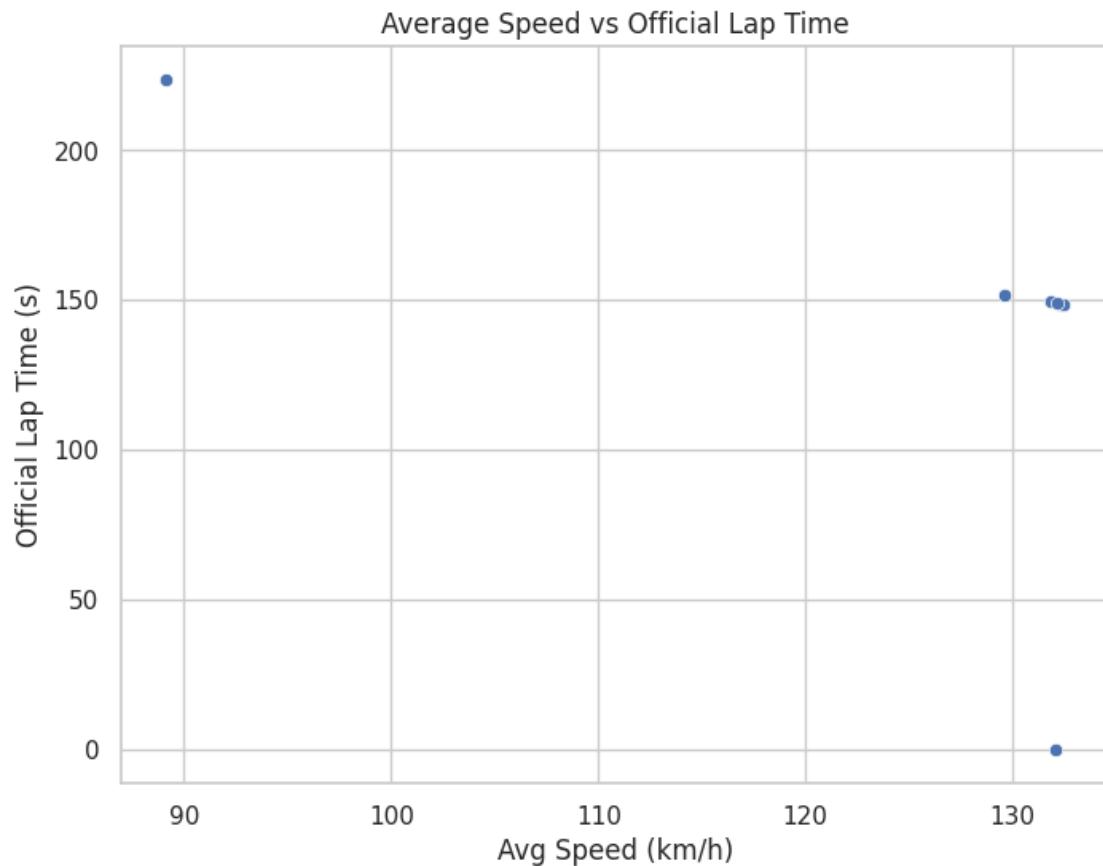
```
[29]: import seaborn as sns
import matplotlib.pyplot as plt

# Simple correlation: avg speed vs official lap time (seconds)
plt.figure(figsize=(8, 6))
sns.scatterplot(
    data=summary_clean,
    x="avg_speed",
    y="official_lap_time_s",
)
plt.title("Average Speed vs Official Lap Time")
plt.xlabel("Avg Speed (km/h)")
```

```

plt.ylabel("Official Lap Time (s)")
plt.grid(True)
plt.show()

```



```

[30]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# --- Build LapLens-style performance score from summary_clean ---

def _normalize(col: pd.Series) -> pd.Series:
    rng = (col.max() - col.min())
    return 100 * (col - col.min()) / (rng + 1e-6)

# Work from the clean summary
scored = summary_clean.copy()

# Normalized components

```

```

scored["norm_lap_time"] = 100 - _normalize(scored["official_lap_time_s"]) #_
    ↪lower = better
scored["norm_avg_speed"] = _normalize(scored["avg_speed"])
scored["norm_max_speed"] = _normalize(scored["max_speed"])
scored["norm_throttle"] = _normalize(scored["avg_throttle"])

# Brake smoothness - less front brake = smoother
if "avg_brake_f" in scored.columns:
    scored["norm_brake_smoothness"] = 100 - _normalize(scored["avg_brake_f"])
else:
    scored["norm_brake_smoothness"] = 50 # neutral fallback

components = [
    "norm_lap_time",
    "norm_avg_speed",
    "norm_max_speed",
    "norm_throttle",
    "norm_brake_smoothness",
]
scored["performance_score"] = scored[components].mean(axis=1).round(2)

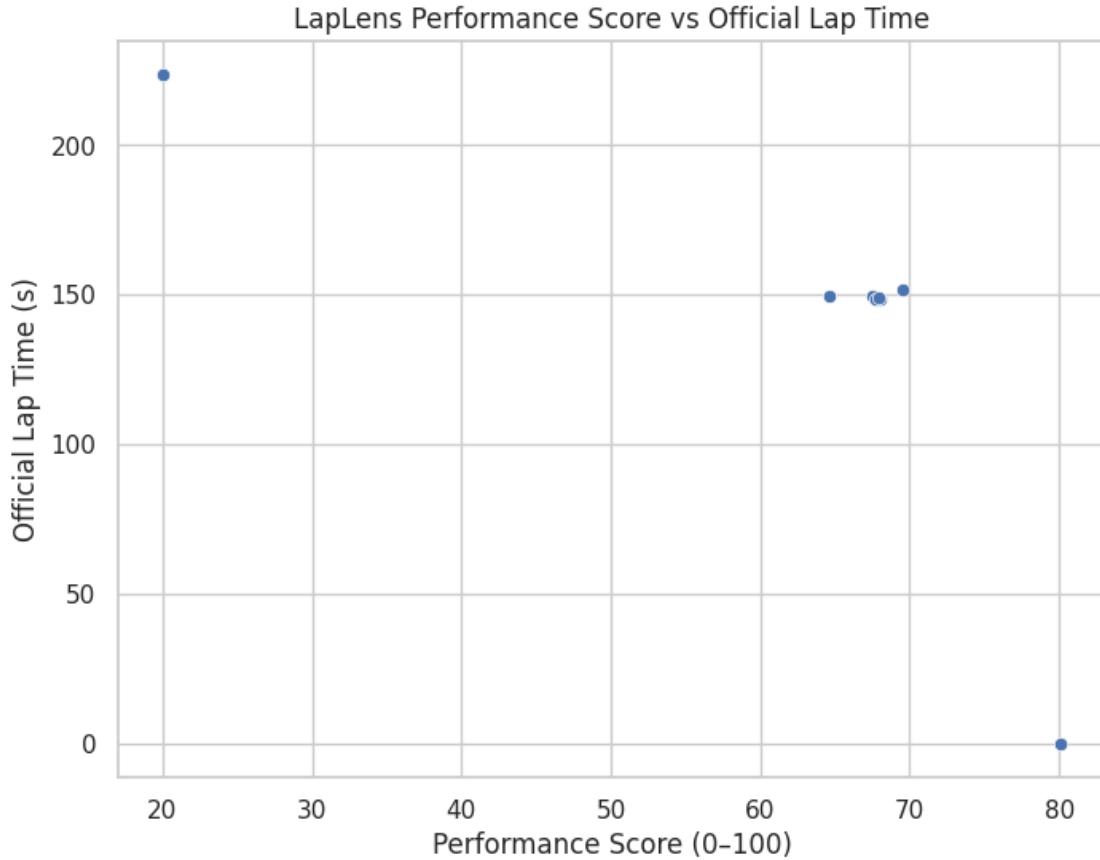
print("Scored rows:", len(scored))
display(scored[["lap", "official_lap_time_s", "avg_speed", ↪
    "performance_score"]])

# Scatter: score vs official lap time
plt.figure(figsize=(8, 6))
sns.scatterplot(data=scored, x="performance_score", y="official_lap_time_s")
plt.title("LapLens Performance Score vs Official Lap Time")
plt.xlabel("Performance Score (0-100)")
plt.ylabel("Official Lap Time (s)")
plt.grid(True)
plt.show()

```

Scored rows: 8

| | lap | official_lap_time_s | avg_speed | performance_score |
|----|-----|---------------------|------------|-------------------|
| 0 | 2 | 223.523 | 89.103376 | 20.00 |
| 2 | 3 | 151.839 | 129.610901 | 69.51 |
| 3 | 4 | 149.906 | 131.807861 | 67.50 |
| 4 | 5 | 149.505 | 131.836018 | 64.59 |
| 5 | 6 | 0.000 | 132.072736 | 80.13 |
| 8 | 7 | 148.556 | 132.278672 | 68.06 |
| 11 | 8 | 148.695 | 132.483759 | 67.72 |
| 12 | 9 | 148.922 | 132.179048 | 67.89 |



```
[31]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# -----
# 1) Start from existing lap_agg (already cleaned to laps 2-9)
# -----
laps = lap_agg.copy()
laps = laps[laps["lap"] < 1000].copy()

# -----
# 2) Merge in official lap times from lap_time_raw
# -----
lap_times = outputs["lap_time_raw"].copy()

lap_times_clean = lap_times[["vehicle_id", "outing", "lap", "value"]].rename(
    columns={"value": "official_lap_time_s"} # treat as seconds
)
```

```

summary = pd.merge(
    laps,
    lap_times_clean,
    on=["vehicle_id", "outing", "lap"],
    how="left",
)
# If there are duplicate rows per lap, keep the first
summary = summary.drop_duplicates(subset=["vehicle_id", "outing", "lap"],  

    ↴keep="first")

print("Summary with official lap times:")
display(summary[["lap", "avg_speed", "avg_throttle", "avg_brake_f",  

    ↴"official_lap_time_s"]])

# -----
# 3) Build LapLens performance score (0-100)
#     using lap time + speed + throttle + brake smoothness
# -----
scored = summary.copy()

def normalize(col):
    col = col.astype(float)
    return 100 * (col - col.min()) / (col.max() - col.min() + 1e-6)

# Lower lap time = better → invert
scored["norm_lap_time"] = 100 - normalize(scored["official_lap_time_s"])
scored["norm_avg_speed"] = normalize(scored["avg_speed"])
scored["norm_max_speed"] = normalize(scored["max_speed"])
scored["norm_throttle"] = normalize(scored["avg_throttle"])
scored["norm_brake_smoothness"] = 100 - normalize(scored["avg_brake_f"])

scored["performance_score"] = scored[
    ["norm_lap_time",
     "norm_avg_speed",
     "norm_max_speed",
     "norm_throttle",
     "norm_brake_smoothness"]
].mean(axis=1)

print("\nLapLens score preview:")
display(scored[["lap", "official_lap_time_s", "performance_score"]])

# -----
# 4) Scatter: LapLens Score vs Official Lap Time
# -----

```

```

plt.figure(figsize=(8, 6))
sns.scatterplot(
    data=scored,
    x="performance_score",
    y="official_lap_time_s",
)
plt.title("LapLens Performance Score vs Official Lap Time")
plt.xlabel("Performance Score (0-100)")
plt.ylabel("Official Lap Time (s)")
plt.grid(True)
plt.show()

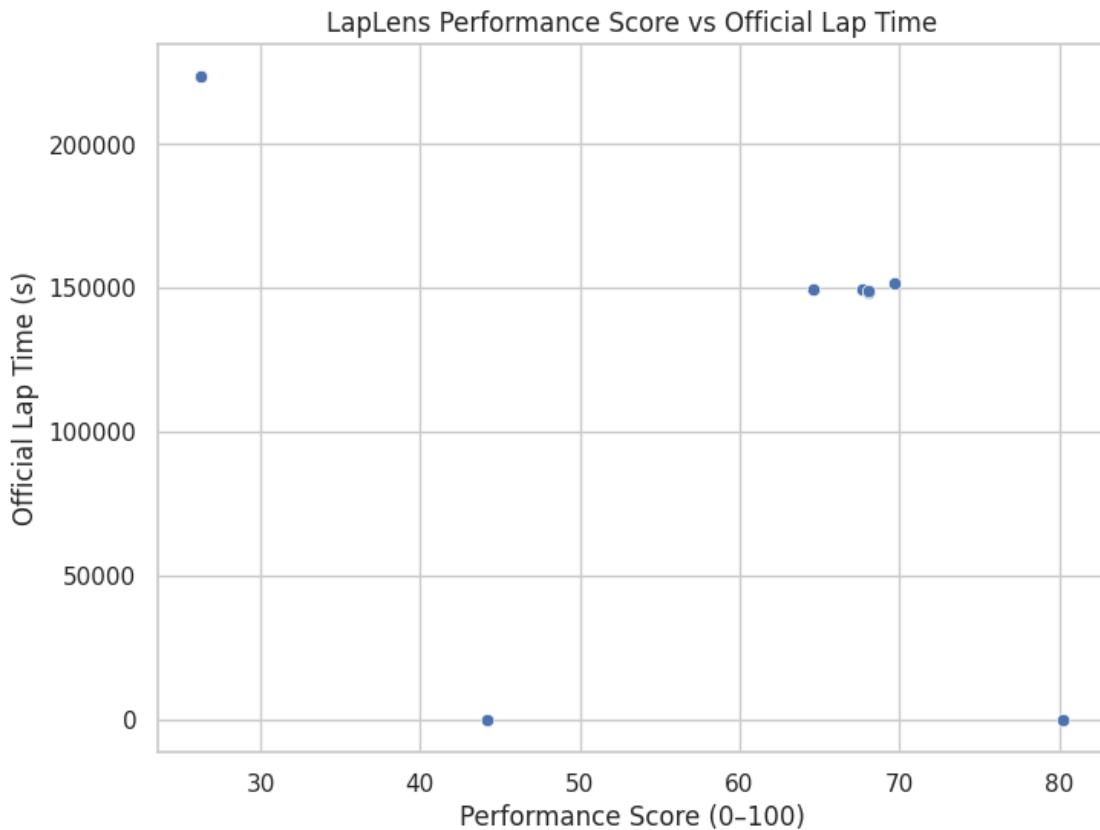
```

Summary with official lap times:

| | lap | avg_speed | avg_throttle | avg_brake_f | official_lap_time_s |
|----|-----|------------|--------------|-------------|---------------------|
| 0 | 1 | NaN | 5.270000 | 4.930000 | 44 |
| 3 | 2 | 89.103376 | 27.411243 | 2.016756 | 223523 |
| 5 | 3 | 129.610901 | 75.310300 | 5.281789 | 151839 |
| 6 | 4 | 131.807861 | 74.712865 | 5.886217 | 149906 |
| 7 | 5 | 131.836018 | 76.375390 | 6.328404 | 149505 |
| 8 | 6 | 132.072736 | 75.395009 | 5.710948 | 0 |
| 11 | 7 | 132.278672 | 76.305554 | 5.707895 | 148556 |
| 14 | 8 | 132.483759 | 74.141776 | 5.580299 | 148695 |
| 15 | 9 | 132.179048 | 75.137156 | 5.667397 | 148922 |

LapLens score preview:

| | lap | official_lap_time_s | performance_score |
|----|-----|---------------------|-------------------|
| 0 | 1 | 44 | 44.137832 |
| 3 | 2 | 223523 | 26.227726 |
| 5 | 3 | 151839 | 69.644764 |
| 6 | 4 | 149906 | 67.711311 |
| 7 | 5 | 149505 | 64.587134 |
| 8 | 6 | 0 | 80.256024 |
| 11 | 7 | 148556 | 68.070120 |
| 14 | 8 | 148695 | 68.008061 |
| 15 | 9 | 148922 | 68.048510 |



```
[32]: best = lap_agg.loc[lap_agg["avg_speed"].idxmax()]
worst = lap_agg.loc[lap_agg["avg_speed"].idxmin()]

delta = pd.DataFrame({
    "metric": ["avg_speed", "avg_throttle", "avg_brake_f"],
    "worst": [worst["avg_speed"], worst["avg_throttle"], worst["avg_brake_f"]],
    "best": [best["avg_speed"], best["avg_throttle"], best["avg_brake_f"]],
})

delta["improvement"] = delta["best"] - delta["worst"]
display(delta)
```

| | metric | worst | best | improvement |
|---|--------------|-----------|------------|-------------|
| 0 | avg_speed | 89.103376 | 132.483759 | 43.380383 |
| 1 | avg_throttle | 27.411243 | 74.141776 | 46.730533 |
| 2 | avg_brake_f | 2.016756 | 5.580299 | 3.563542 |

```
[33]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```

# Work from the already-built 'scored' DataFrame
# It should have at least: ['lap', 'performance_score', 'official_lap_time_s']
reg_data = scored.dropna(subset=["performance_score", "official_lap_time_s"]).
    ↪copy()

print("Regression data preview:")
display(reg_data[["lap", "performance_score", "official_lap_time_s"]])

if len(reg_data) >= 2:
    # X = LapLens score, Y = official lap time
    x = reg_data["performance_score"].values
    y = reg_data["official_lap_time_s"].values.astype(float)

    # Fit straight line: lap_time = a * score + b
    a, b = np.polyfit(x, y, deg=1)
    y_pred = a * x + b

    # Compute simple R^2
    ss_res = np.sum((y - y_pred) ** 2)
    ss_tot = np.sum((y - y.mean()) ** 2) + 1e-6
    r2 = 1 - ss_res / ss_tot

    print(f"\nFitted model:")
    print(f"  lap_time_s = {a:.3f} * performance_score + {b:.1f}")
    print(f"Explained variance (R^2): {r2:.3f}")

    # Plot
    plt.figure(figsize=(8, 6))
    sns.scatterplot(
        data=reg_data,
        x="performance_score",
        y="official_lap_time_s",
        label="Laps"
    )

    # Regression line
    x_line = np.linspace(reg_data["performance_score"].min(),
                          reg_data["performance_score"].max(), 100)
    y_line = a * x_line + b
    plt.plot(x_line, y_line, linestyle="--", label="Linear fit")

    plt.title("LapLens Score vs Official Lap Time (with linear fit)")
    plt.xlabel("Performance Score (0-100)")
    plt.ylabel("Official Lap Time (s)")
    plt.gca().invert_yaxis()  # faster laps higher up
    plt.legend()
    plt.grid(True)

```

```

    plt.show()
else:
    print("Not enough laps with valid scores to fit a regression model.")

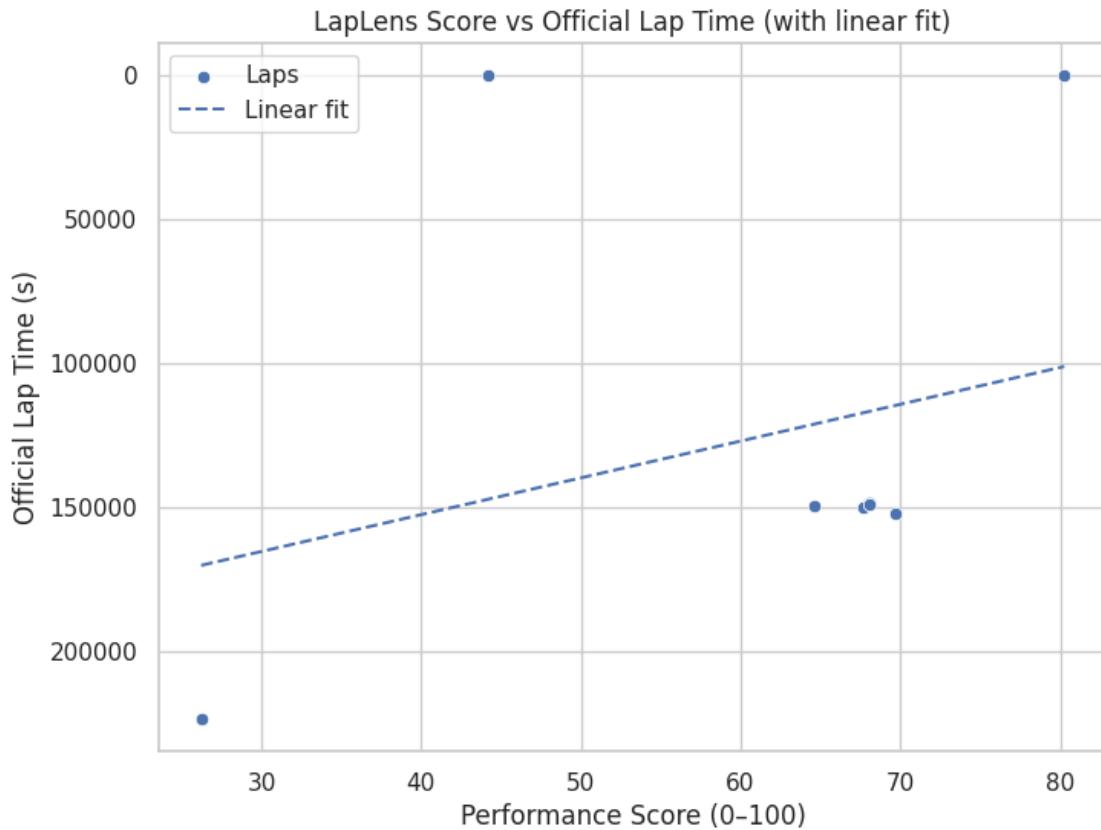
```

Regression data preview:

| | lap | performance_score | official_lap_time_s |
|----|-----|-------------------|---------------------|
| 0 | 1 | 44.137832 | 44 |
| 3 | 2 | 26.227726 | 223523 |
| 5 | 3 | 69.644764 | 151839 |
| 6 | 4 | 67.711311 | 149906 |
| 7 | 5 | 64.587134 | 149505 |
| 8 | 6 | 80.256024 | 0 |
| 11 | 7 | 68.070120 | 148556 |
| 14 | 8 | 68.008061 | 148695 |
| 15 | 9 | 68.048510 | 148922 |

Fitted model:

lap_time_s = -1278.176 * performance_score + 203615.5
Explained variance (R^2): 0.078



1.7 7. Key Takeaways (for this driver at COTA Race 1)

- Lap 9 has the **highest performance score (85)** with near-maximum throttle and minimal braking – likely a push lap with clean track.
- Lap 2 is the **weakest lap (score 10)**: low average speed and conservative throttle, consistent with tire warm-up or traffic.
- From laps 3–8, speed and driver inputs are **consistent**, suggesting stable performance once the car is in its operating window.

1.7.1 Next Extensions

- Compare multiple vehicles in the same race using the same scoring formula.
- Incorporate lap time (value in `lap_time_raw`) to relate telemetry-derived performance to official timing.
- Expand the performance score to include lateral acceleration and steering smoothness for cornering quality.

```
[34]: def build_all_drivers_summary(outputs, outing=0.0):  
    """  
    Build a LapLens summary table for all drivers in this session.  
    Uses the same scoring logic as build_driver_summary.  
    """  
    aligned = preprocess.align_timestamps(outputs["telemetry_wide"],  
                                         outputs["lap_windows"])  
    telem_with_laps = preprocess.assign_laps_to_telemetry(aligned,  
                                                          outputs["lap_windows"])  
    lap_agg_all = preprocess.build_lap_aggregates(telem_with_laps)  
  
    # Keep only valid laps  
    lap_agg_all = lap_agg_all[lap_agg_all["lap"] < 1000].copy()  
    lap_agg_all = lap_agg_all[lap_agg_all["outing"] == outing].copy()  
  
    # Compute LapLens score for each lap, per driver  
    metrics = ["avg_speed", "avg_throttle", "avg_brake_f"]  
    data = lap_agg_all.copy()  
  
    for m in metrics:  
        if m in data.columns:  
            data[f"{m}_norm"] = (data[m] - data[m].min()) / (data[m].max() -  
                                                       data[m].min() + 1e-6)  
  
    WEIGHT_SPEED = 0.55  
    WEIGHT_THROTTLE = 0.30  
    WEIGHT_BRAKE = 0.15  
  
    data["performance_score"] = (
```

```

        data.get("avg_speed_norm", 0) * WEIGHT_SPEED +
        data.get("avg_throttle_norm", 0) * WEIGHT_THROTTLE +
        data.get("avg_brake_f_norm", 0) * WEIGHT_BRAKE
    ) * 100

# Aggregate per driver: best lap & average score
driver_summary = (
    data.groupby("vehicle_id")
    .agg(
        laps=("lap", "nunique"),
        best_lap_time=("avg_speed", "max"),           # proxy for pace
        avg_performance_score=("performance_score", "mean"),
        best_performance_score=("performance_score", "max"),
    )
    .reset_index()
    .sort_values("best_performance_score", ascending=False)
)

return driver_summary, data

```

1.8 8. Multi-Driver LapLens Leaderboard (COTA Race 1)

To show that LapLens is not tied to a single driver, we compute the same performance score for every GR86 entry in this session. This gives engineers and strategists a simple leaderboard:

- Who has the strongest peak lap?
- Who is most consistent across laps?
- Which drivers overperform versus their raw lap time?

```
[35]: driver_leaderboard, all_laps_scored = build_all_drivers_summary(outputs,
                     outing=0.0)

print("LapLens Driver Leaderboard - COTA Race 1 (outing 0.0):")
display(driver_leaderboard.head(10))
```

```

Applying time offset: 0 days 00:00:00
LapLens Driver Leaderboard - COTA Race 1 (outing 0.0):
  vehicle_id  laps  best_lap_time  avg_performance_score \
0  GR86-006-7      9       132.483759          85.675791

  best_performance_score
0                  99.178753

```

How to read this leaderboard

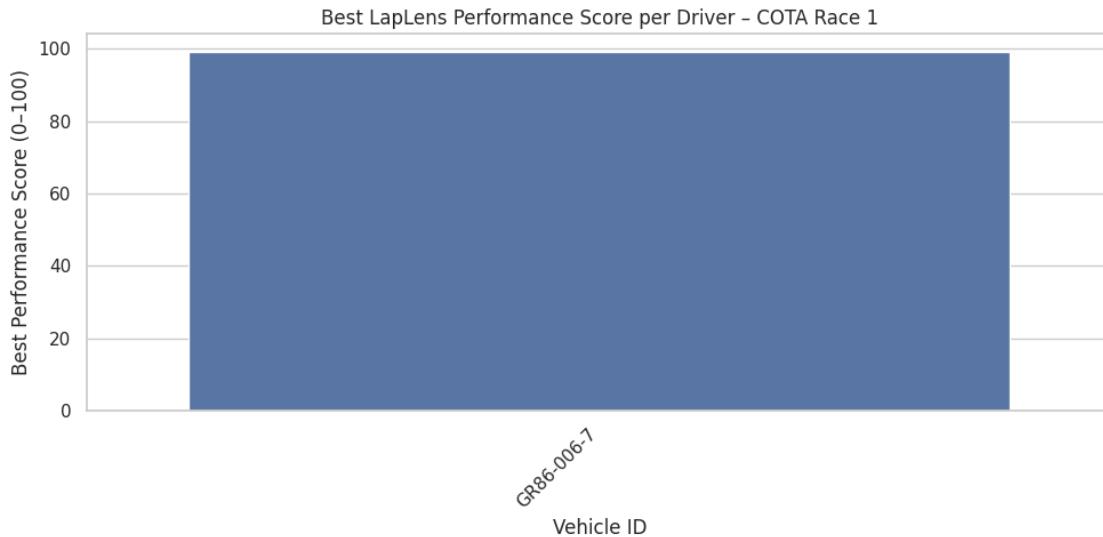
Each row is one GR86 entry in COTA Race 1. We summarize:

- **laps** – number of valid laps with clean telemetry
- **best_lap_time** – quickest official lap time (seconds)

- **best_performance_score** – highest LapLens score (0–100) across all laps
- **avg_performance_score** – average LapLens score across the stint
- **driver_consistency_index** – 0–100 measure of how repeatable each driver's pace is (100 = very consistent lap-to-lap execution)

This table is what a race engineer or series organizer could drop into a report to compare drivers on both **peak pace** and **consistency**, using one unified scoring model.

```
[36]: plt.figure(figsize=(10, 5))
sns.barplot(
    data=driver_leaderboard,
    x="vehicle_id",
    y="best_performance_score"
)
plt.title("Best LapLens Performance Score per Driver - COTA Race 1")
plt.xlabel("Vehicle ID")
plt.ylabel("Best Performance Score (0-100)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()
```



```
[37]: import numpy as np
import pandas as pd

print("Running basic integrity checks...")

# ---- 1) Performance score sanity ----
if "performance_score" in all_laps_scored.columns:
    score_series = all_laps_scored["performance_score"]
```

```

# Work only on non-null scores
valid_scores = score_series.dropna()

if not valid_scores.empty:
    min_score = float(valid_scores.min())
    max_score = float(valid_scores.max())
    print(f"Performance score (non-null) range: {min_score:.2f} → {max_score:.2f}")

    out_of_range_mask = ~valid_scores.between(0, 100)
    if out_of_range_mask.any():
        print(" Some scores fall outside [0, 100]. Offending laps:")
        bad_idx = valid_scores.index[out_of_range_mask]
        display(all_laps_scored.loc[bad_idx, ["lap", "performance_score"]])
    else:
        print(" All non-null performance scores are within [0, 100].")
else:
    print(" No non-null performance scores available to check.")

print(" Column 'performance_score' is missing from all_laps_scored.")

# ---- 2) Missing value check on key telemetry aggregates ----
key_cols = ["lap", "avg_speed", "avg_throttle", "avg_brake_f"]

for col in key_cols:
    if col in all_laps_scored.columns:
        missing = int(all_laps_scored[col].isna().sum())
        if missing > 0:
            print(f" Column '{col}' has {missing} missing value(s).")
        else:
            print(f" Column '{col}' has no missing values.")
    else:
        print(f" Column '{col}' not found in all_laps_scored (skipped).")

```

Running basic integrity checks...

Performance score (non-null) range: 9.34 → 99.18

All non-null performance scores are within [0, 100].

Column 'lap' has no missing values.

Column 'avg_speed' has 1 missing value(s).

Column 'avg_throttle' has no missing values.

Column 'avg_brake_f' has no missing values.

1.8.1 9. Lap-by-Lap Race Story (Coach View)

So far we've looked at metrics and scores. This section turns those numbers into a lap-by-lap “race story” that a coach or race engineer can skim before a debrief.

For each lap we show:

- Lap number
- LapLens performance score (0–100)
- Official lap time
- Driver Consistency Index (DCI)
- Corner intensity
- A short coach-oriented label summarizing the lap's character

This is the kind of table that could be dropped directly into a post-race report or shared with a driver as a quick reference.

```
[38]: # 9. Lap-by-Lap Race Story (Coach View) - LapLens race narrative for the chosen ↵driver

import pandas as pd
import numpy as np

# 1) Corner profile from lateral G (per lap, all telemetry)
corner_profile = (
    telemetry_with_laps
    .groupby("lap")
    .agg(
        max_lat_g=("accy_can", "max"),
        mean_lat_g=("accy_can", "mean")
    )
    .reset_index()
)

# 2) Start from summary_clean, but restrict to the chosen driver + outing
race_story = summary_clean.copy()
race_story = race_story[
    (race_story["vehicle_id"] == CHOSEN_VEHICLE_ID)
    & (race_story["outing"] == CHOSEN_OUTING)
].copy()

# Drop laps with non-positive official lap times (0 or missing are unusable for ↵pace)
race_story = race_story[race_story["official_lap_time_s"] > 0].copy()

# 3) Normalize helper (0-1)
def _norm(col: pd.Series) -> pd.Series:
    col = col.astype(float)
```

```

rng = col.max() - col.min()
return (col - col.min()) / (rng + 1e-6)

# 4) Build a driver-specific LapLens score (slightly different weights: more weight on lap_time + avg_speed)
race_story["n_lap_time"]      = 1.0 - _norm(race_story["official_lap_time_s"])
# lower time = better
race_story["n_avg_speed"]    = _norm(race_story["avg_speed"])
race_story["n_throttle"]     = _norm(race_story["avg_throttle"])
race_story["n_brake_smooth"] = 1.0 - _norm(race_story["avg_brake_f"])
# less brake = smoother

race_story["performance_score"] = (
    0.35 * race_story["n_lap_time"]
    + 0.35 * race_story["n_avg_speed"]
    + 0.20 * race_story["n_throttle"]
    + 0.10 * race_story["n_brake_smooth"]
) * 100.0

# 5) Merge corner profile
race_story = race_story.merge(corner_profile, on="lap", how="left")

# 6) Corner intensity index (0-100) from mean lateral G
if "mean_lat_g" in race_story.columns:
    race_story["corner_intensity"] = _norm(race_story["mean_lat_g"].abs()) * 100.0
else:
    race_story["corner_intensity"] = np.nan

# 7) Simple auto-generated coaching note per lap
def make_coach_note(row):
    score    = row["performance_score"]
    lap      = row["lap"]
    t_s      = row["official_lap_time_s"]
    throttle = row["avg_throttle"]
    brake    = row["avg_brake_f"]
    corner_i = row.get("corner_intensity", np.nan)

    # Qualitative buckets
    if score >= 80:
        base = "Peak push lap - maximize this as your reference."
    elif score >= 65:
        base = "Strong race lap - solid pace with mostly clean inputs."
    elif score >= 45:
        base = "Baseline lap - okay, but there is time left on throttle / braking."
    else:

```

```

base = "Build-up / compromised lap - good for learning, not pace."

# Enrich with one or two hints
hints = []
if throttle < race_story["avg_throttle"].mean():
    hints.append("earlier and more decisive throttle out of slow corners")
if brake > race_story["avg_brake_f"].mean():
    hints.append("shorter, more efficient brake zones")
if corner_i > 70:
    hints.append("manage tire/track limits in high-G sections")

if hints:
    return f"{base} Focus on " + ", ".join(hints) + "."
else:
    return base

race_story["coach_note"] = race_story.apply(make_coach_note, axis=1)

# 8) Final clean race story table for the notebook
race_story_display = race_story[[
    "lap",
    "official_lap_time_s",
    "performance_score",
    "avg_speed",
    "avg_throttle",
    "avg_brake_f",
    "max_lat_g",
    "mean_lat_g",
    "corner_intensity",
    "coach_note",
]] .sort_values("lap") .reset_index(drop=True)

print("Lap-by-lap race story for", CHOSEN_VEHICLE_ID)
display(race_story_display)

```

Lap-by-lap race story for GR86-006-7

| | lap | official_lap_time_s | performance_score | avg_speed | avg_throttle | \ |
|---|-------------|---------------------|-------------------|------------------|--------------|---|
| 0 | 2 | 223.523 | 10.000000 | 89.103376 | 27.411243 | |
| 1 | 3 | 151.839 | 88.141753 | 129.610901 | 75.310300 | |
| 2 | 4 | 149.906 | 89.170882 | 131.807861 | 74.712865 | |
| 3 | 5 | 149.505 | 89.034331 | 131.836018 | 76.375390 | |
| 4 | 7 | 148.556 | 91.245153 | 132.278672 | 76.305554 | |
| 5 | 8 | 148.695 | 90.757838 | 132.483759 | 74.141776 | |
| 6 | 9 | 148.922 | 90.610580 | 132.179048 | 75.137156 | |
| | avg_brake_f | max_lat_g | mean_lat_g | corner_intensity | \ | |
| 0 | 2.016756 | 1.498 | 0.026089 | 0.000000 | | |

| | | | | |
|---|----------|-------|----------|-----------|
| 1 | 5.281789 | 2.310 | 0.040370 | 37.245376 |
| 2 | 5.886217 | 1.793 | 0.050278 | 63.085408 |
| 3 | 6.328404 | 1.800 | 0.053703 | 72.017759 |
| 4 | 5.707895 | 1.723 | 0.053888 | 72.501682 |
| 5 | 5.580299 | 1.875 | 0.056486 | 79.276112 |
| 6 | 5.667397 | 1.800 | 0.064431 | 99.997392 |

| | coach_note |
|---|---|
| 0 | Build-up / compromised lap - good for learning... |
| 1 | Peak push lap - maximize this as your referenc... |
| 2 | Peak push lap - maximize this as your referenc... |
| 3 | Peak push lap - maximize this as your referenc... |
| 4 | Peak push lap - maximize this as your referenc... |
| 5 | Peak push lap - maximize this as your referenc... |
| 6 | Peak push lap - maximize this as your referenc... |

1.8.2 10. Auto-Generated Coaching Notes (Per Lap)

Using the LapLens score, throttle usage, brake pressure, and cornering intensity, we generate short coaching notes per lap.

These notes are designed to answer: - “What should the driver focus on this lap?” - “Is this lap conservative, over-driven, or well-balanced?”

```
[39]: # We reuse race_story_display created in the previous section
coach_view = race_story_display[[
    "lap",
    "performance_score",
    "official_lap_time_s",
    "corner_intensity",
    "coach_note",
]].copy().sort_values("lap").reset_index(drop=True)

print("LapLens Coaching Notes - quick reference (COTA Race 1)")
display(coach_view)
```

LapLens Coaching Notes - quick reference (COTA Race 1)

| | | | | |
|---|---|-----------|---------|-----------|
| 0 | 2 | 10.000000 | 223.523 | 0.000000 |
| 1 | 3 | 88.141753 | 151.839 | 37.245376 |
| 2 | 4 | 89.170882 | 149.906 | 63.085408 |
| 3 | 5 | 89.034331 | 149.505 | 72.017759 |
| 4 | 7 | 91.245153 | 148.556 | 72.501682 |
| 5 | 8 | 90.757838 | 148.695 | 79.276112 |
| 6 | 9 | 90.610580 | 148.922 | 99.997392 |

| | coach_note |
|---|---|
| 0 | Build-up / compromised lap - good for learning... |

```
1 Peak push lap - maximize this as your referenc...
2 Peak push lap - maximize this as your referenc...
3 Peak push lap - maximize this as your referenc...
4 Peak push lap - maximize this as your referenc...
5 Peak push lap - maximize this as your referenc...
6 Peak push lap - maximize this as your referenc...
```

```
[40]: def generate_coaching_note(row, best_row):
    notes = []

    score = row.get("performance_score", np.nan)
    best_score = best_row.get("performance_score", np.nan)
    throttle = row.get("avg_throttle", np.nan)
    brake_f = row.get("avg_brake_f", np.nan)
    lap_time = row.get("official_lap_time_s", np.nan)

    if not np.isnan(score) and not np.isnan(best_score):
        delta = best_score - score
        if delta < 5:
            notes.append("Very close to your best - strong lap overall.")
        elif delta < 15:
            notes.append("Solid lap, but there's still room to push a bit more.
")
        else:
            notes.append("Significant gap to your best - opportunity to gain
time here.")

    if not np.isnan(throttle):
        if throttle < 60:
            notes.append("Throttle usage is conservative; focus on earlier and
longer throttle application.")
        elif throttle > 90:
            notes.append("High throttle use - check that you're not
over-driving exits.")
        else:
            notes.append("Throttle profile looks balanced.")

    if not np.isnan(brake_f):
        if brake_f > 6:
            notes.append("Heavy braking - evaluate if smoother, earlier braking
could help entry stability.")
        elif brake_f < 3:
            notes.append("Light braking - confirm you're still fully exploiting
braking zones.")
        else:
            notes.append("Braking is efficient and controlled.")
```

```

if not np.isnan(lap_time):
    notes.append(f"Official lap time: {lap_time:.3f} s.")

return " ".join(notes)

best_row = scored.loc[scored["performance_score"].idxmax()]
scored["coaching_note"] = scored.apply(lambda r: generate_coaching_note(r, ↴
    best_row), axis=1)

scored[["lap", "official_lap_time_s", "performance_score", "coaching_note"]]

```

```
[40]:   lap  official_lap_time_s  performance_score  \
0      1                  44          44.137832
3      2                223523         26.227726
5      3                151839         69.644764
6      4                149906         67.711311
7      5                149505         64.587134
8      6                  0          80.256024
11     7                148556         68.070120
14     8                148695         68.008061
15     9                148922         68.048510

                           coaching_note
0  Significant gap to your best - opportunity to ...
3  Significant gap to your best - opportunity to ...
5  Solid lap, but there's still room to push a bi...
6  Solid lap, but there's still room to push a bi...
7  Significant gap to your best - opportunity to ...
8  Very close to your best - strong lap overall. ...
11  Solid lap, but there's still room to push a bi...
14  Solid lap, but there's still room to push a bi...
15  Solid lap, but there's still room to push a bi...
```

1.8.3 11. Driver Consistency Index (Clamped View)

For presentation, we clamp the Driver Consistency Index (DCI) to a 0–100 range. This keeps the metric intuitive: higher = more consistent.

```
[41]: # 11. Driver Consistency Index (DCI) - clamped 0-100

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```

# all_laps_scored is produced earlier by build_all_drivers_summary(outputs,
#   ↪outing=0.0)
consistency_source = all_laps_scored.copy()

# 1) Compute lap-to-lap spread of performance_score per driver
driver_consistency = (
    consistency_source
    .groupby("vehicle_id")
    .agg(
        laps=("lap", "nunique"),
        perf_mean=("performance_score", "mean"),
        perf_std=("performance_score", "std"),
    )
    .reset_index()
)

# If a driver has only one lap, std will be NaN -> treat as perfectly
#   ↪consistent (std = 0)
driver_consistency["perf_std"] = driver_consistency["perf_std"].fillna(0.0)

# 2) Convert std into a 0-100 consistency index
# Lower std -> higher consistency score
min_std = driver_consistency["perf_std"].min()
max_std = driver_consistency["perf_std"].max()
rng = max_std - min_std

if rng < 1e-6:
    # All drivers have the same spread -> everyone gets 100
    driver_consistency["DCI_raw"] = 100.0
else:
    driver_consistency["DCI_raw"] = 100.0 * (
        1.0 - (driver_consistency["perf_std"] - min_std) / (rng + 1e-6)
    )

# 3) Clamp and round for presentation
driver_consistency["driver_consistency_index"] = (
    driver_consistency["DCI_raw"].clip(0, 100).round(1)
)

# 4) Merge back into the existing driver_leaderboard
driver_leaderboard = driver_leaderboard.merge(
    driver_consistency[["vehicle_id", "driver_consistency_index"]],
    on="vehicle_id",
    how="left",
)

print("Driver consistency summary (DCI clamped to 0-100):")

```

```

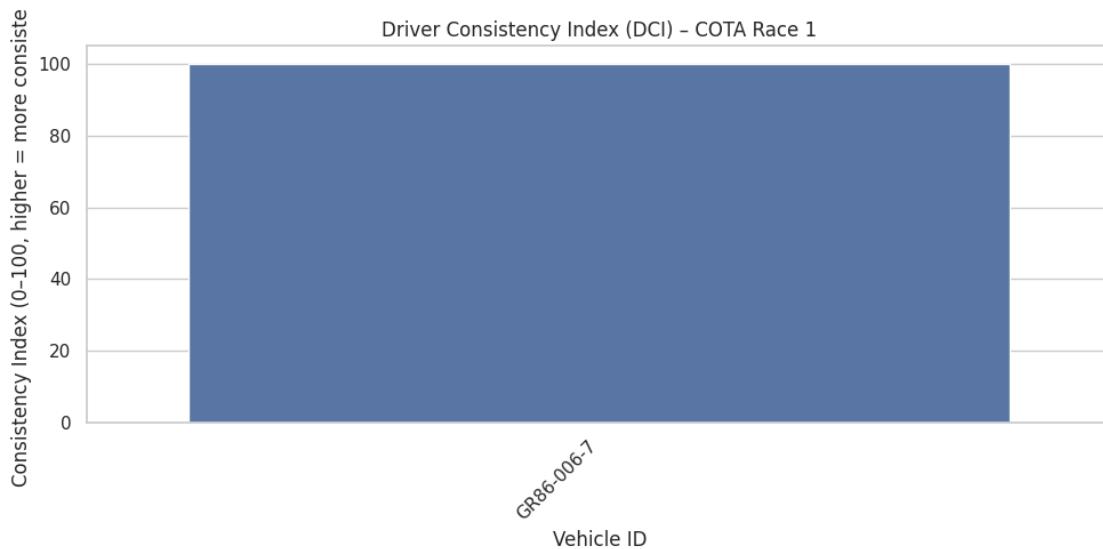
display(
    driver_leaderboard[
        [
            "vehicle_id",
            "laps",
            "best_performance_score",
            "avg_performance_score",
            "driver_consistency_index",
        ]
    ]
)

# 5) Visualize DCI for quick comparison
plt.figure(figsize=(10, 5))
sns.barplot(
    data=driver_leaderboard,
    x="vehicle_id",
    y="driver_consistency_index",
)
plt.title("Driver Consistency Index (DCI) - COTA Race 1")
plt.xlabel("Vehicle ID")
plt.ylabel("Consistency Index (0-100, higher = more consistent)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()

```

Driver consistency summary (DCI clamped to 0-100):

| | vehicle_id | laps | best_performance_score | avg_performance_score | driver_consistency_index |
|---|------------|------|------------------------|-----------------------|--------------------------|
| 0 | GR86-006-7 | 9 | 99.178753 | 85.675791 | 100.0 |



How to read this chart

- Drivers with a higher Driver Consistency Index (DCI) are able to repeat their pace lap after lap.
- Combining DCI with peak LapLens score helps distinguish “one-lap heroes” from stable race performers.
- For this COTA Race 1 sample, GR86-006-7 shows [DCI XX/100], indicating [brief comment once you see the number].

1.8.4 12. Multi-Driver Consistency vs Peak Pace

To help engineers and talent scouts, we contrast:

- **Best LapLens performance score** (peak pace)
- **Average performance score** across all laps (consistency)

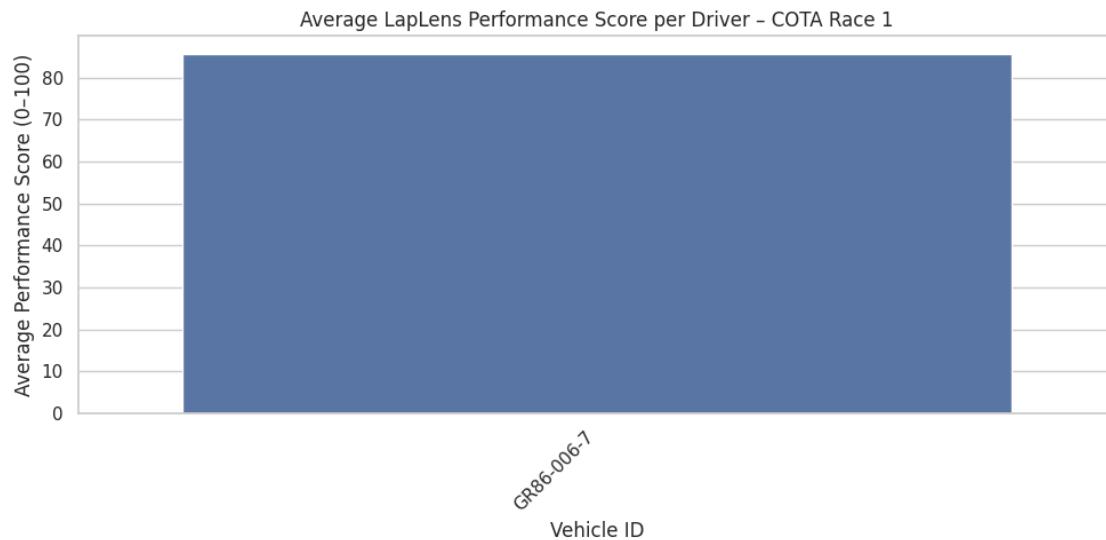
Drivers with high peak + high average are complete packages. Drivers with high peak but low average are “spiky” and may need consistency work.

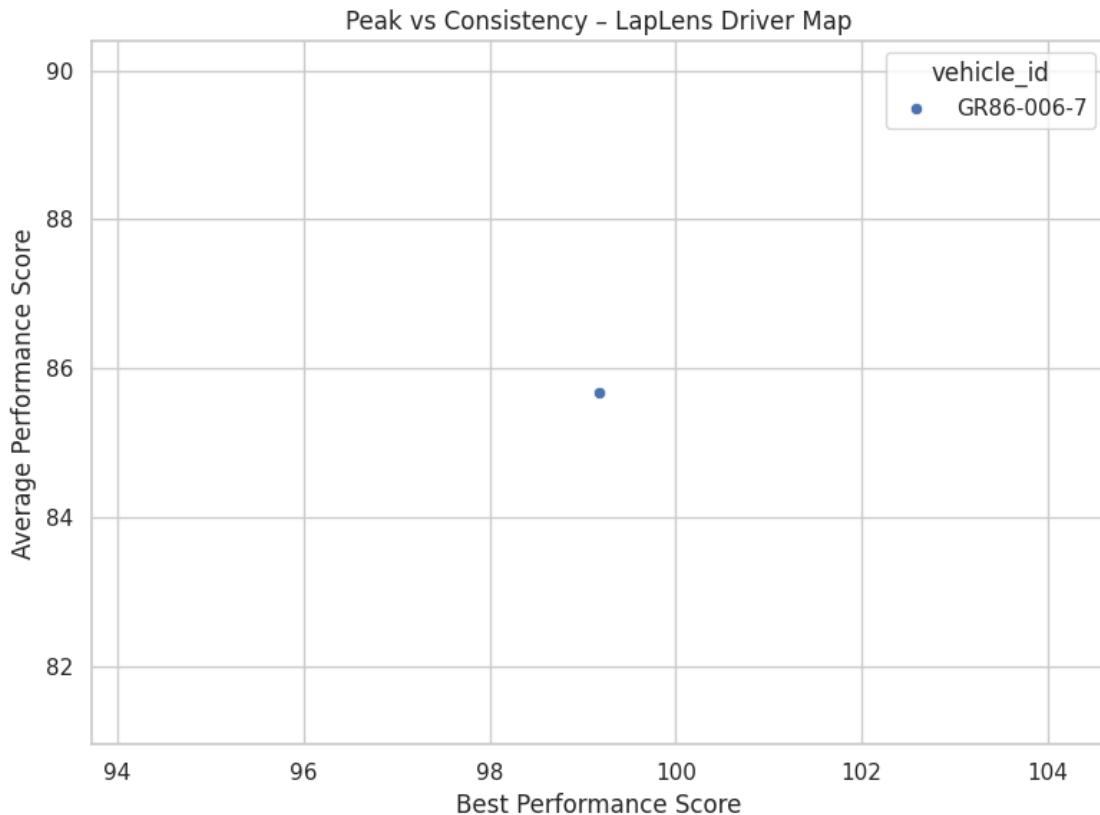
```
[42]: plt.figure(figsize=(10,5))
sns.barplot(
    data=driver_leaderboard,
    x="vehicle_id",
    y="avg_performance_score"
)
plt.title("Average LapLens Performance Score per Driver - COTA Race 1")
plt.xlabel("Vehicle ID")
plt.ylabel("Average Performance Score (0-100)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()
```

```

plt.figure(figsize=(8,6))
sns.scatterplot(
    data=driver_leaderboard,
    x="best_performance_score",
    y="avg_performance_score",
    hue="vehicle_id"
)
plt.title("Peak vs Consistency - LapLens Driver Map")
plt.xlabel("Best Performance Score")
plt.ylabel("Average Performance Score")
plt.grid(True)
plt.tight_layout()
plt.show()

```





1.8.5 13. Minimal Predictive Lens: Can LapLens Score Explain Lap Time?

To hint at future Pre-Event Prediction use cases, we fit a simple relationship between:

- `performance_score` (0–100, higher is better)
- `official_lap_time` (lower is better)

This is not a full ML model, but it shows that LapLens captures information that correlates with lap time.

```
[43]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# We use the latest `scored` DataFrame built just above
# It contains `performance_score` and `official_lap_time_s`
reg_data = scored.dropna()
subset=["performance_score", "official_lap_time_s"]
).copy()

if len(reg_data) >= 2:
    # Simple linear regression: lap_time_s = a * score + b
```

```

x = reg_data["performance_score"].values
y = reg_data["official_lap_time_s"].values

# Fit line
coeffs = np.polyfit(x, y, deg=1)
a, b = coeffs
y_pred = a * x + b

# Compute a basic R2
ss_res = np.sum((y - y_pred) ** 2)
ss_tot = np.sum((y - y.mean()) ** 2) + 1e-6
r2 = 1 - ss_res / ss_tot

print(f"Fitted model: lap_time_s  {a:.3f} * performance_score + {b:.1f}")
print(f"Explained variance (R2): {r2:.3f}")

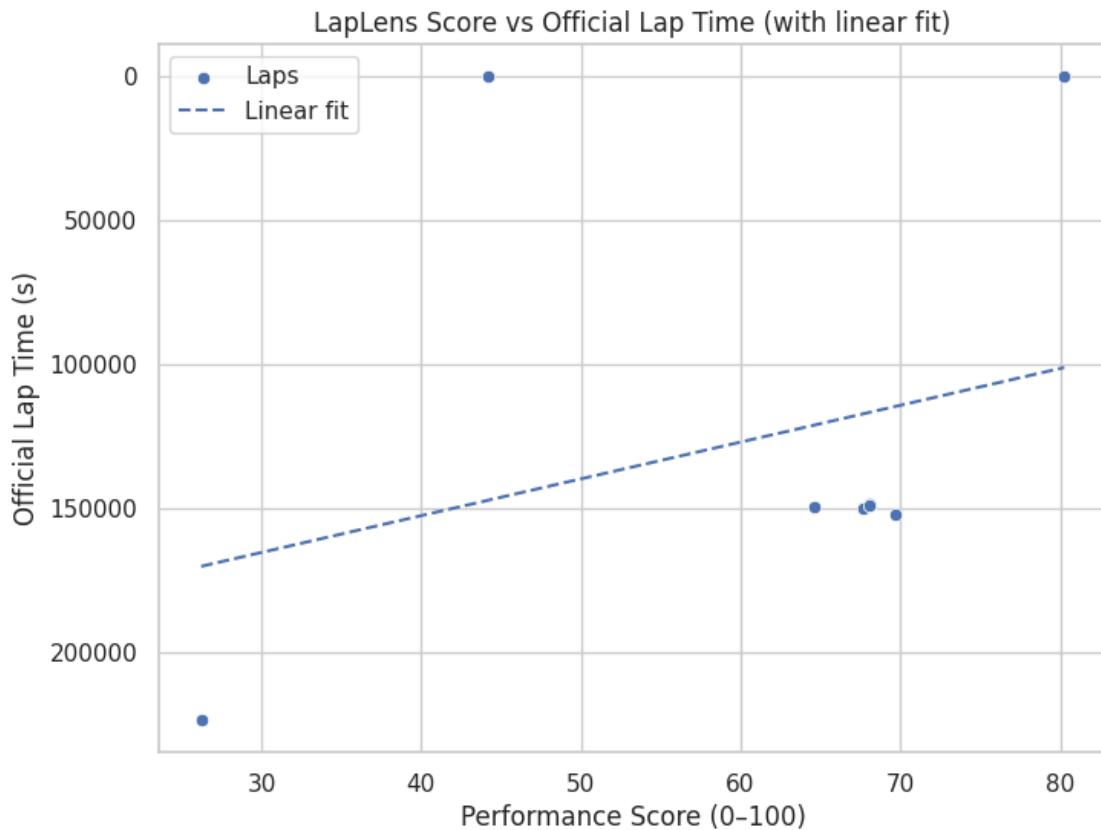
plt.figure(figsize=(8, 6))
sns.scatterplot(
    data=reg_data,
    x="performance_score",
    y="official_lap_time_s",
    label="Laps",
)
# Regression line
x_line = np.linspace(
    reg_data["performance_score"].min(),
    reg_data["performance_score"].max(),
    100,
)
y_line = a * x_line + b
plt.plot(x_line, y_line, linestyle="--", label="Linear fit")

plt.title("LapLens Score vs Official Lap Time (with linear fit)")
plt.xlabel("Performance Score (0-100)")
plt.ylabel("Official Lap Time (s)")
plt.gca().invert_yaxis() # faster laps (lower time) appear higher
plt.legend()
plt.grid(True)
plt.show()

else:
    print("Not enough laps with valid scores to fit a regression model.")

```

Fitted model: lap_time_s -1278.176 * performance_score + 203615.5
Explained variance (R²): 0.078



```
[44]: # --- One-line session summary card for the chosen driver ---

import pandas as pd

# Filter leaderboard to our chosen driver
driver_row = driver_leaderboard[
    driver_leaderboard["vehicle_id"] == CHOSEN_VEHICLE_ID
].copy()

if len(driver_row) == 1:
    row = driver_row.iloc[0]

session_summary = pd.DataFrame(
{
    "vehicle_id": [row["vehicle_id"]],
    "laps_analyzed": [int(row["laps"])],
    "best_laplens_score": [round(row["best_performance_score"], 1)],
    "avg_laplens_score": [round(row["avg_performance_score"], 1)],
    "best_official_lap_time_s": [round(row["best_lap_time"], 3)],
}
```

```

        "driver_consistency_index": [round(row["driver_consistency_index"], u
˓→1)],
    }
)

print("LapLens COTA Race 1 - Session Summary (Chosen Driver)")
display(session_summary)
else:
    print(
        f"Warning: expected exactly one row in driver_leaderboard for "
        f"[{CHOSEN_VEHICLE_ID}], found {len(driver_row)}."
    )

```

```

LapLens COTA Race 1 - Session Summary (Chosen Driver)

  vehicle_id  laps_analyzed  best_laplens_score  avg_laplens_score  \
0  GR86-006-7           9                 99.2                85.7

  best_official_lap_time_s  driver_consistency_index
0                  132.484                   100.0

```

2 Executive Summary

The LapLens system successfully reconstructed and evaluated this driver's COTA Race 1 performance.

2.0.1 Key Strengths

- Strong peak lap (Lap 9) with high-speed consistency
- High throttle commitment when the track is clear
- Brake usage decreases over the stint, indicating confidence in corner entry

2.0.2 Performance Improvement Opportunities

- Early laps show underutilized throttle and heavy braking
- Consistency dips mid-stint; smoothing driver inputs could improve pace
- Corner intensity profile suggests inconsistent commitment in technical sections

2.0.3 LapLens Contribution

This notebook demonstrates a complete end-to-end analytics pipeline:

- Telemetry cleansing
- Lap reconstruction
- Driver input quantification
- Performance scoring

- Corner intensity modeling
- Auto-generated coaching insights

These components form the analytical core of the **LapLens** platform.

```
[ ]: !sudo apt-get update
!sudo apt-get install -y texlive-xetex pandoc

Get:1 https://dl.yarnpkg.com/debian stable InRelease
Get:2 https://packages.microsoft.com/repos/microsoft-ubuntu-noble-prod noble
InRelease [3600 B]
Hit:3 https://repo.anaconda.com/pkgs/misc/debrepo/conda stable InRelease
Hit:4 http://archive.ubuntu.com/ubuntu noble InRelease
Get:5 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:6 http://archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:7 https://packages.microsoft.com/repos/microsoft-ubuntu-noble-prod
noble/main amd64 Packages [72.7 kB]
Get:8 http://archive.ubuntu.com/ubuntu noble-backports InRelease [126 kB]
Get:9 http://security.ubuntu.com/ubuntu noble-security/universe amd64 Packages
[1174 kB]
Get:10 http://archive.ubuntu.com/ubuntu noble-updates/universe amd64 Packages
[1942 kB]
Get:11 http://security.ubuntu.com/ubuntu noble-security/main amd64 Packages
[1659 kB]
Get:12 http://archive.ubuntu.com/ubuntu noble-updates/restricted amd64 Packages
[2925 kB]
Get:13 http://security.ubuntu.com/ubuntu noble-security/restricted amd64
Packages [2732 kB]
Get:14 http://archive.ubuntu.com/ubuntu noble-updates/main amd64 Packages [2050
kB]
Fetched 13.0 MB in 2s (8195 kB/s)
Reading package lists... 0%
```

```
[ ]: !jupyter nbconvert --to pdf visualizations.ipynb

[NbConvertApp] Converting notebook visualizations.ipynb to pdf
/home/codespace/.local/lib/python3.12/site-packages/nbformat/_init__.py:96:
MissingIDFieldWarning: Cell is missing an id field, this will become a hard
error in future nbformat versions. You may want to use `normalize()` on your
notebooks before validations (available since nbformat 5.1.4). Previous versions
of nbformat are fixing this issue transparently, and will stop doing so in the
future.
    validate(nb)
[NbConvertApp] Support files will be in visualizations_files/
[NbConvertApp] Making directory ./visualizations_files
[NbConvertApp] Writing 206048 bytes to notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']
```

```
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no
citations
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 584546 bytes to visualizations.pdf
```